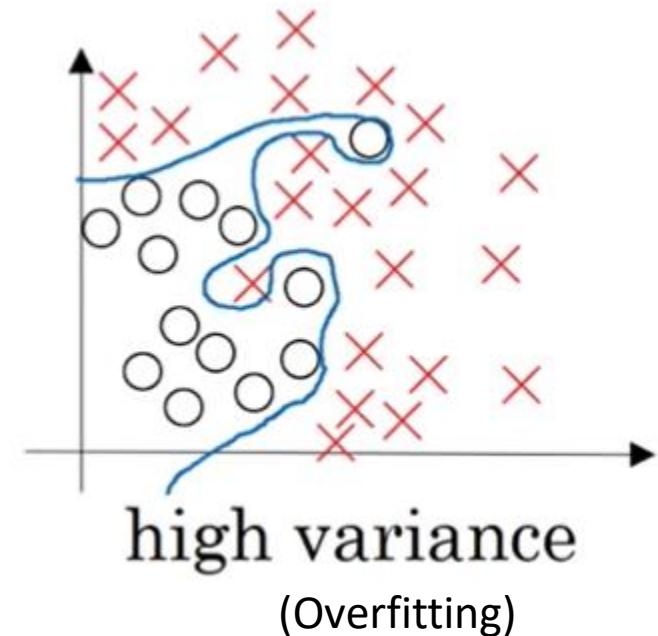
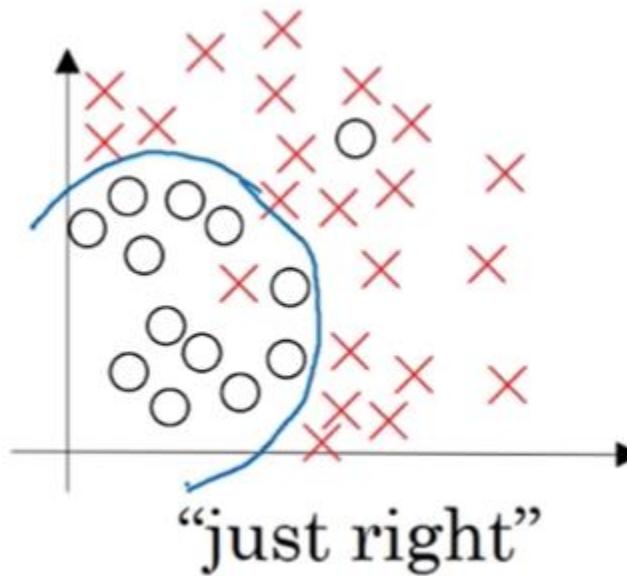
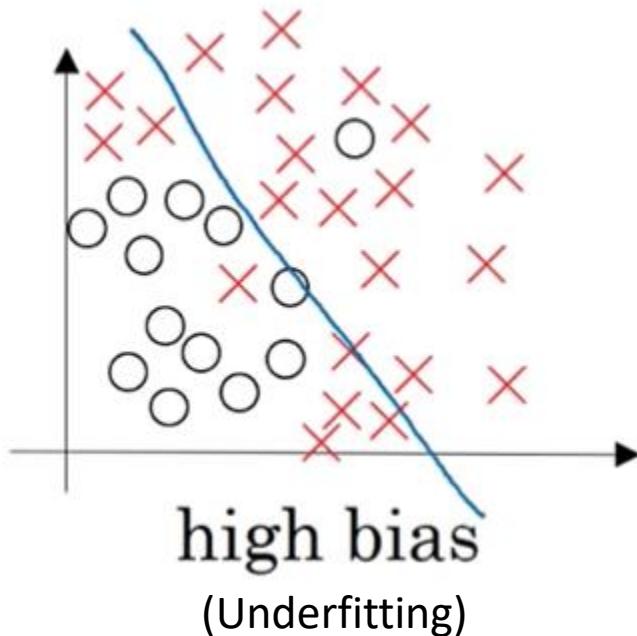


Regularization

Chandranath Adak

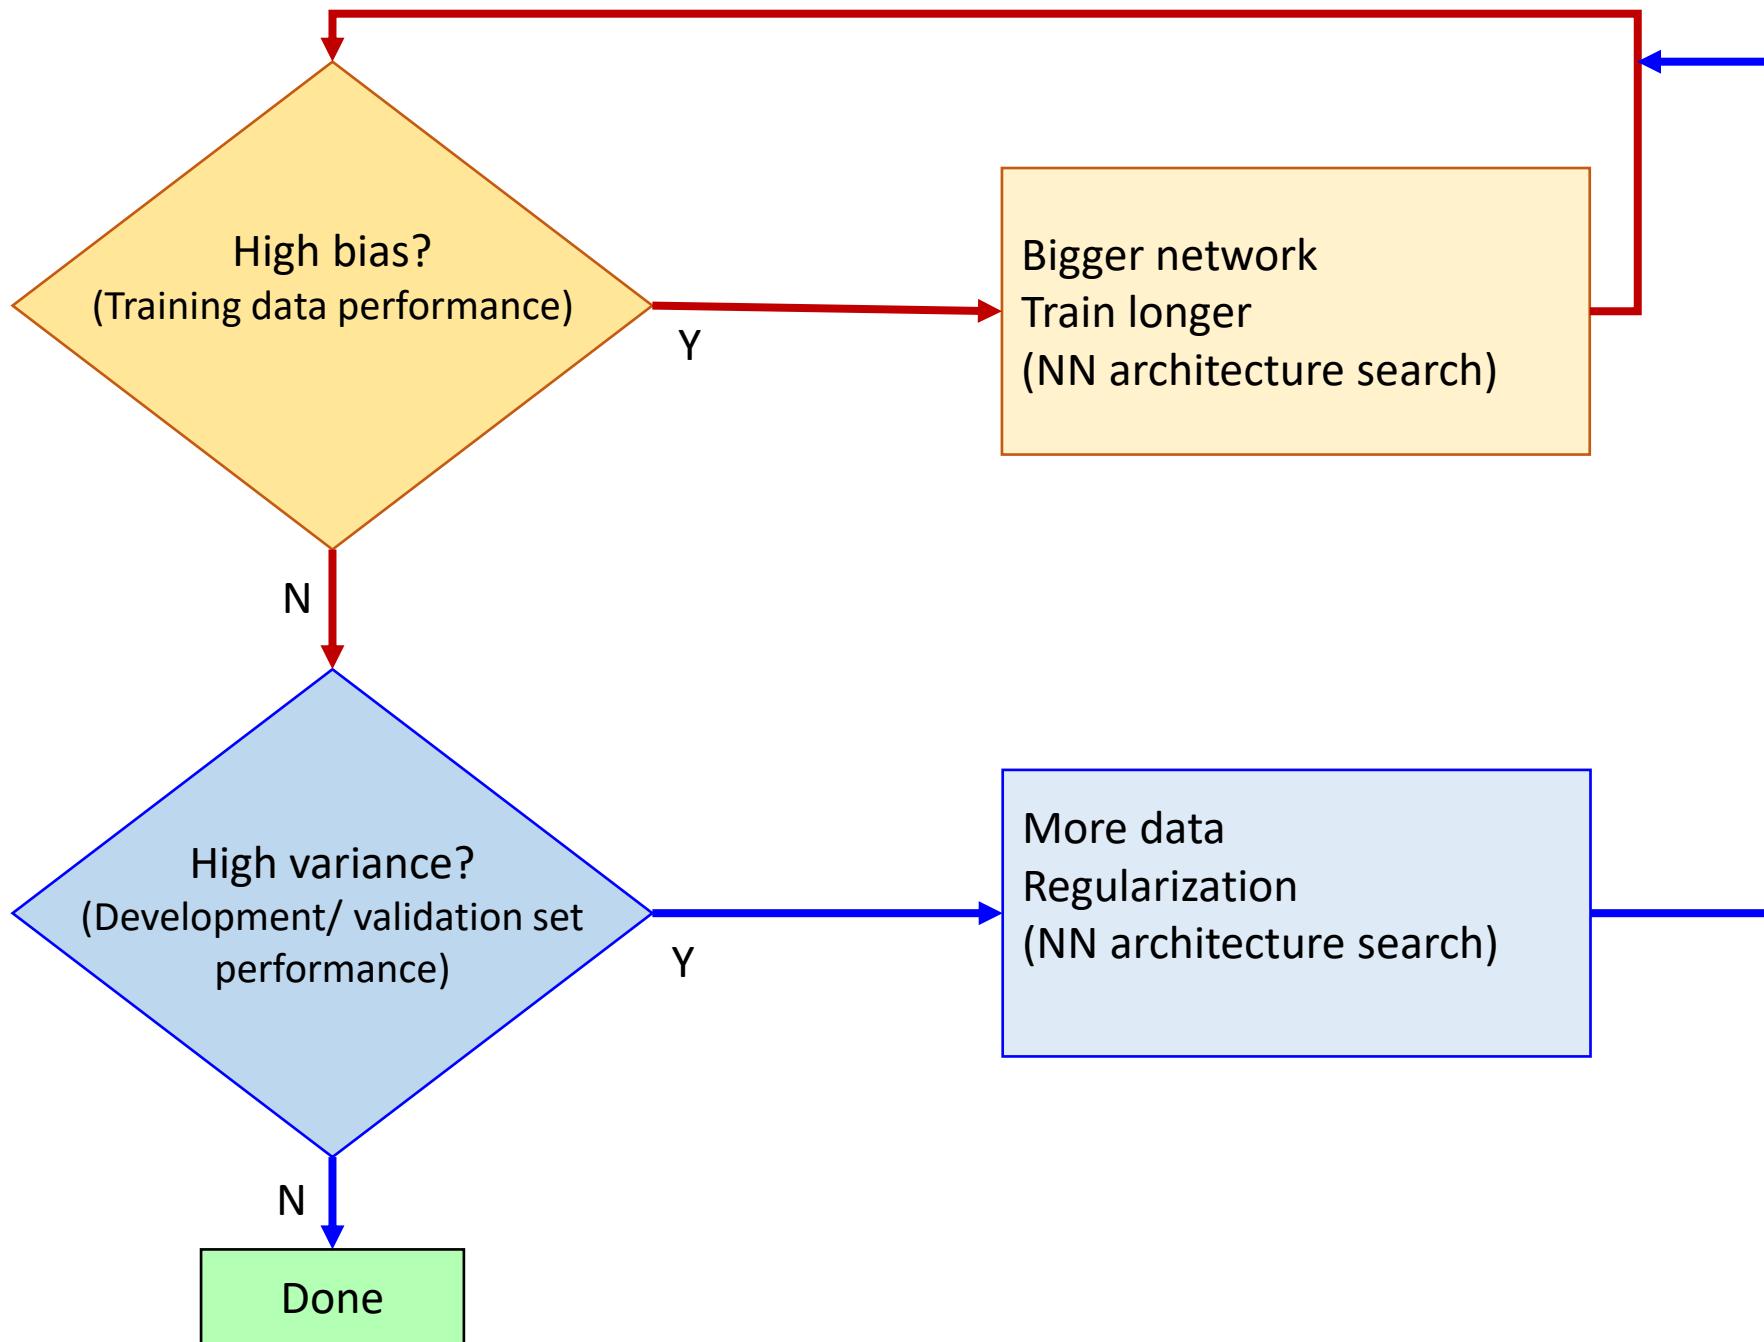
Dept. of CSE, IIT Patna

High Bias vs. High Variance



Training error	1%	15%	15%	0.5%
Testing error	11%	16%	30%	1%
Infer	High variance	High bias	High bias & High variance	Low bias & Low variance

Assumption: Human/ optimal (Bayes') error $\approx 0\%$



Regularization

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \|w\|^2 + \frac{\lambda}{2m} \|b\|^2$$

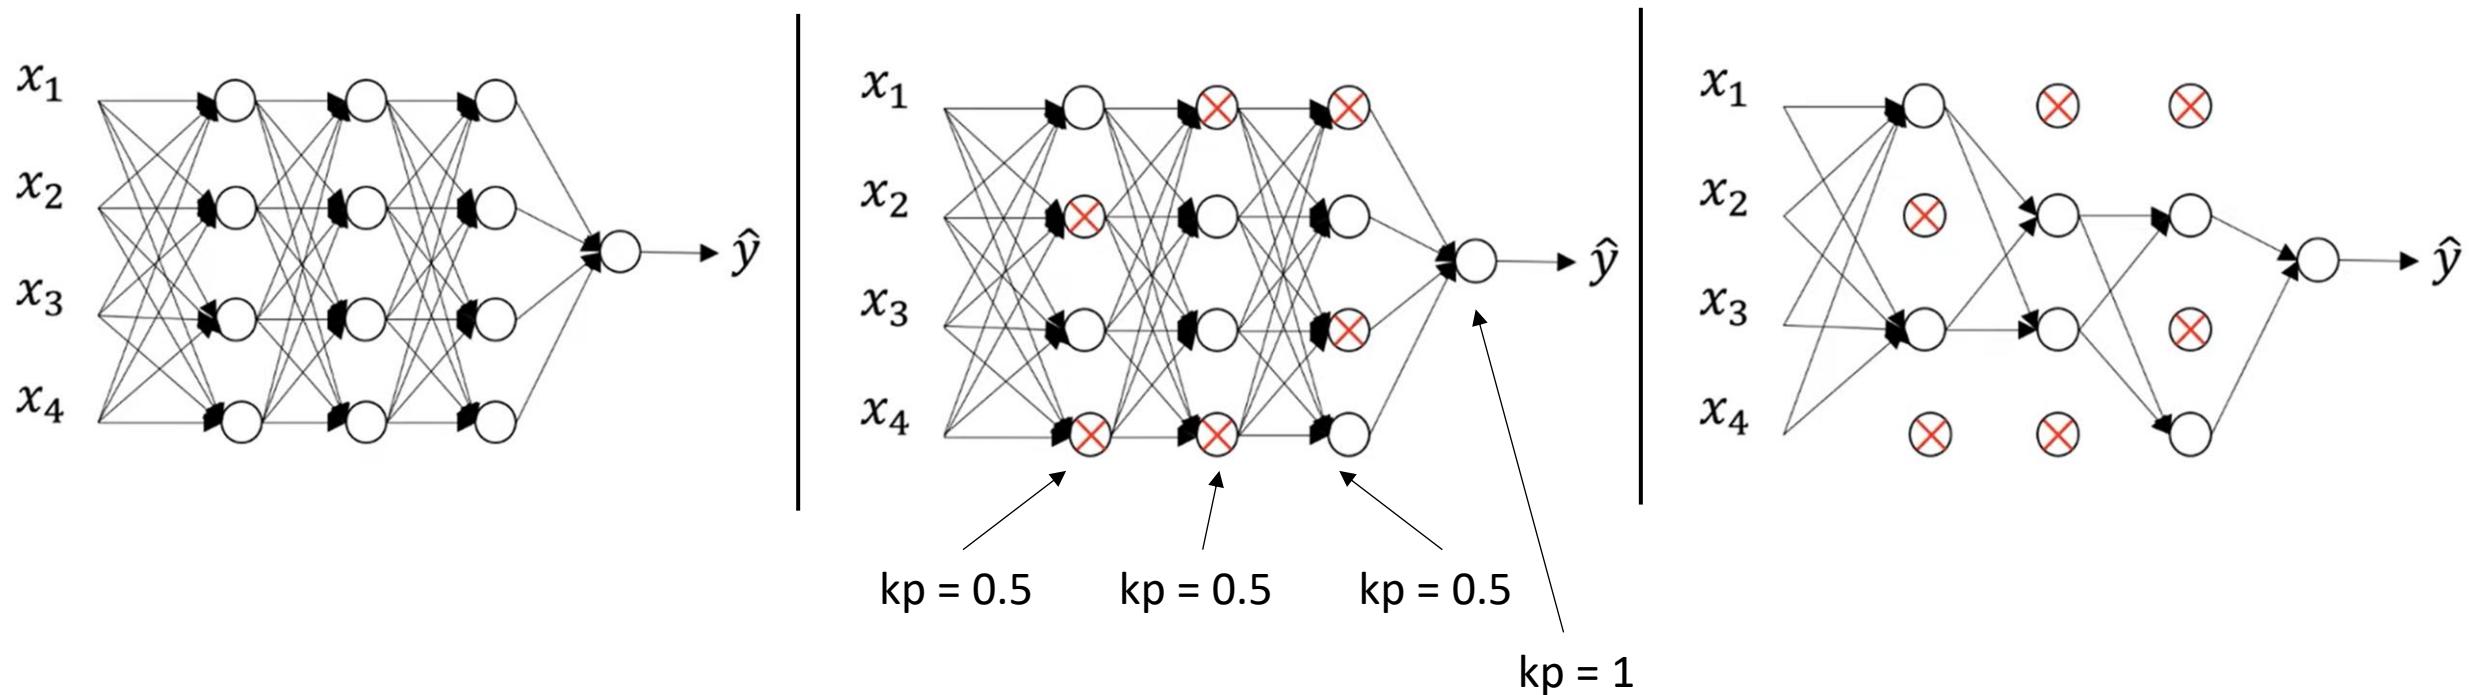
$$\text{L2 regularization} = \|w\|^2 = \sum_{j=1}^{n_x} w_j^2 = w^T w$$

$$\text{L1 regularization} = \|w\|_1 = \sum_{j=1}^{n_x} |w_j|$$

λ increases: underfitting

λ decreases: overfitting

Dropout Regularization



keep_prob (kp) = 1 → No Dropout

Intuition: Can't rely on any one set of features, so have to spread-out weights

Data Augmentation



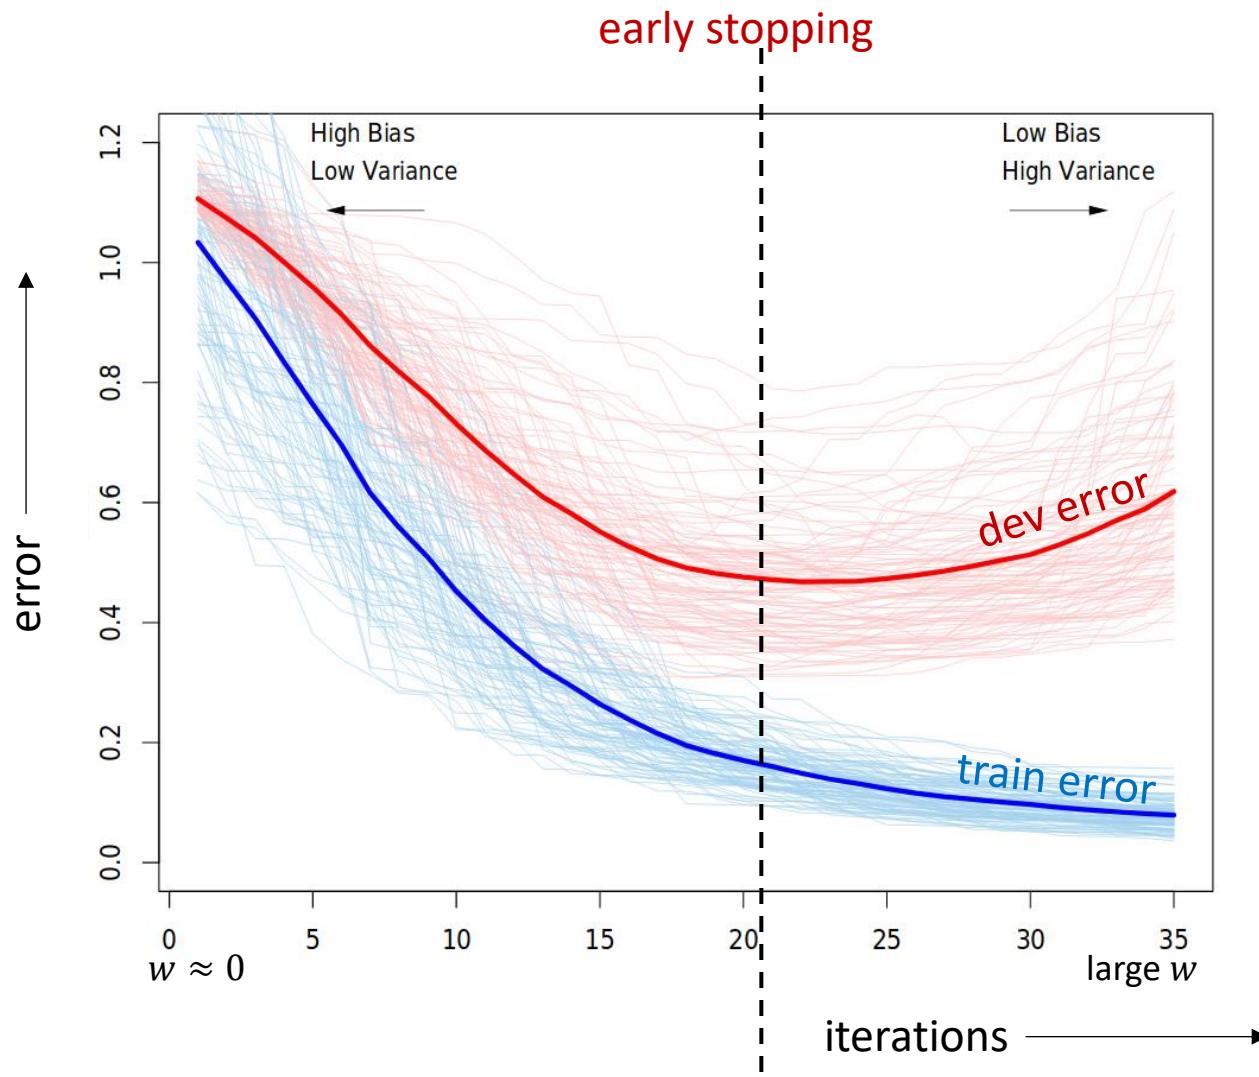
4



Data Augmentation



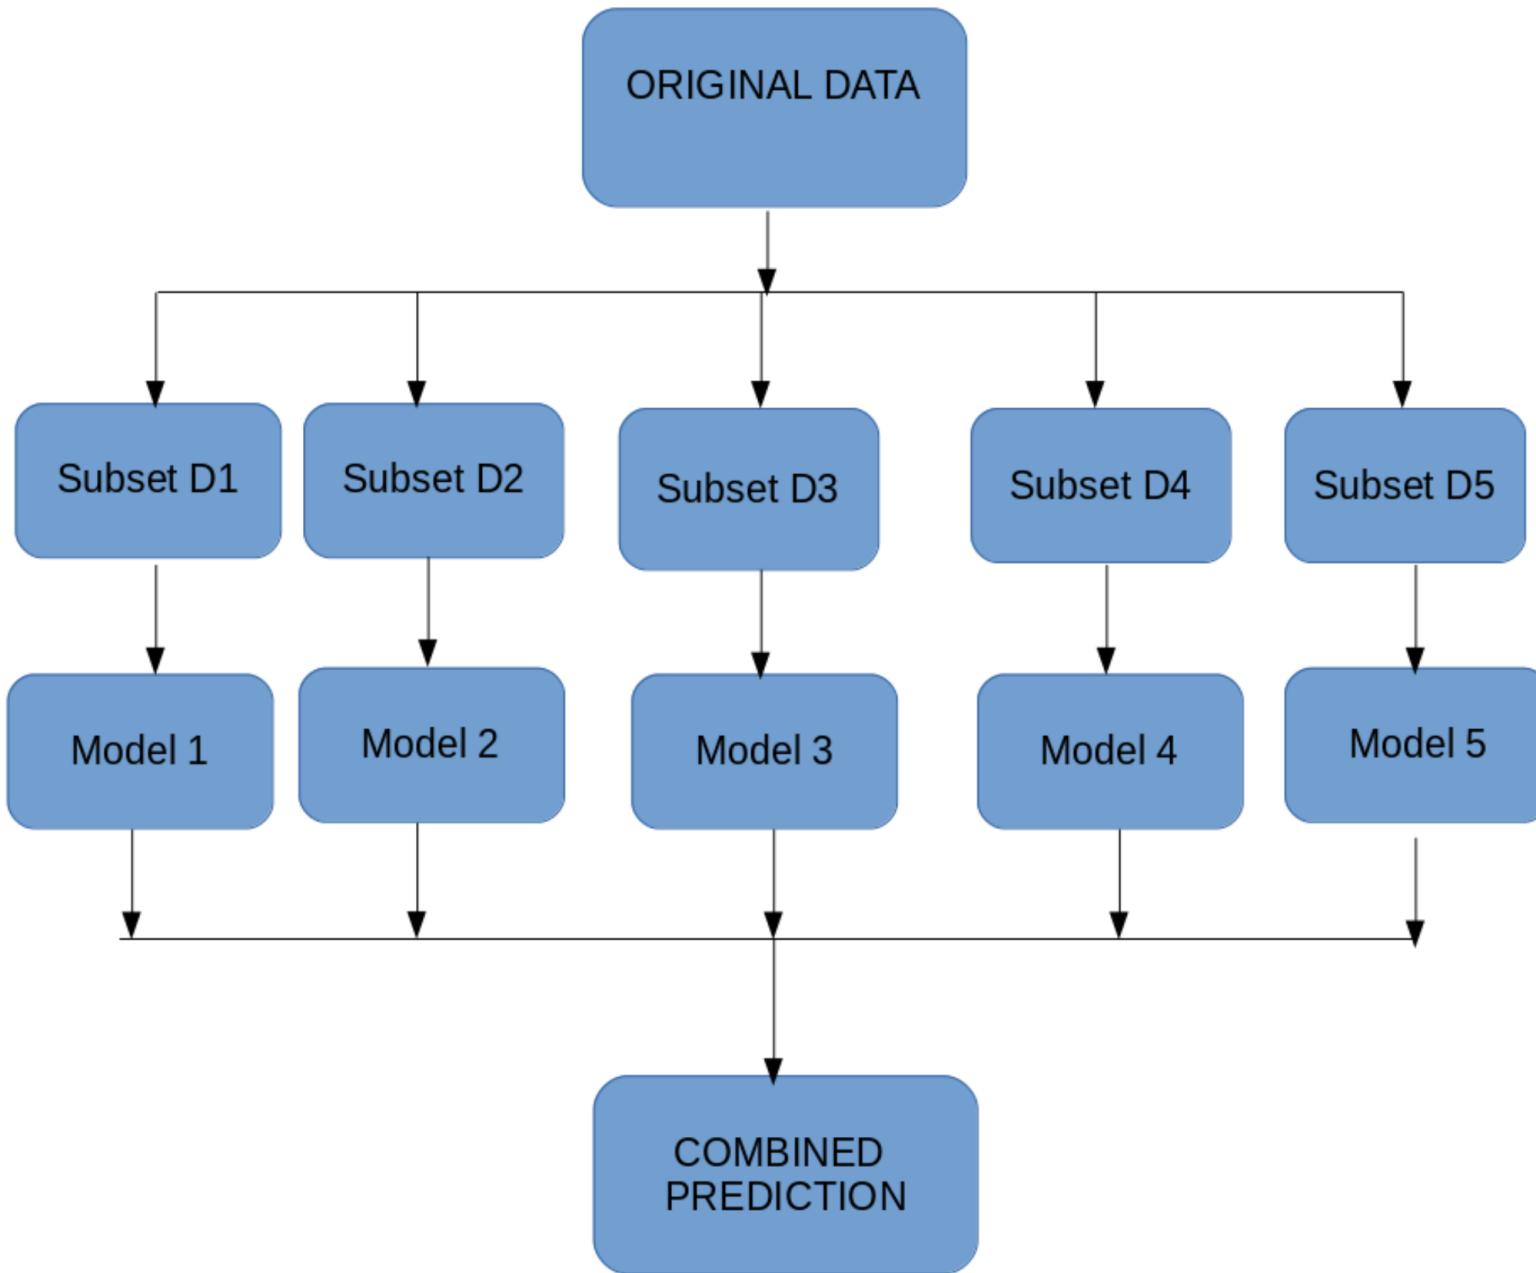
Early Stopping



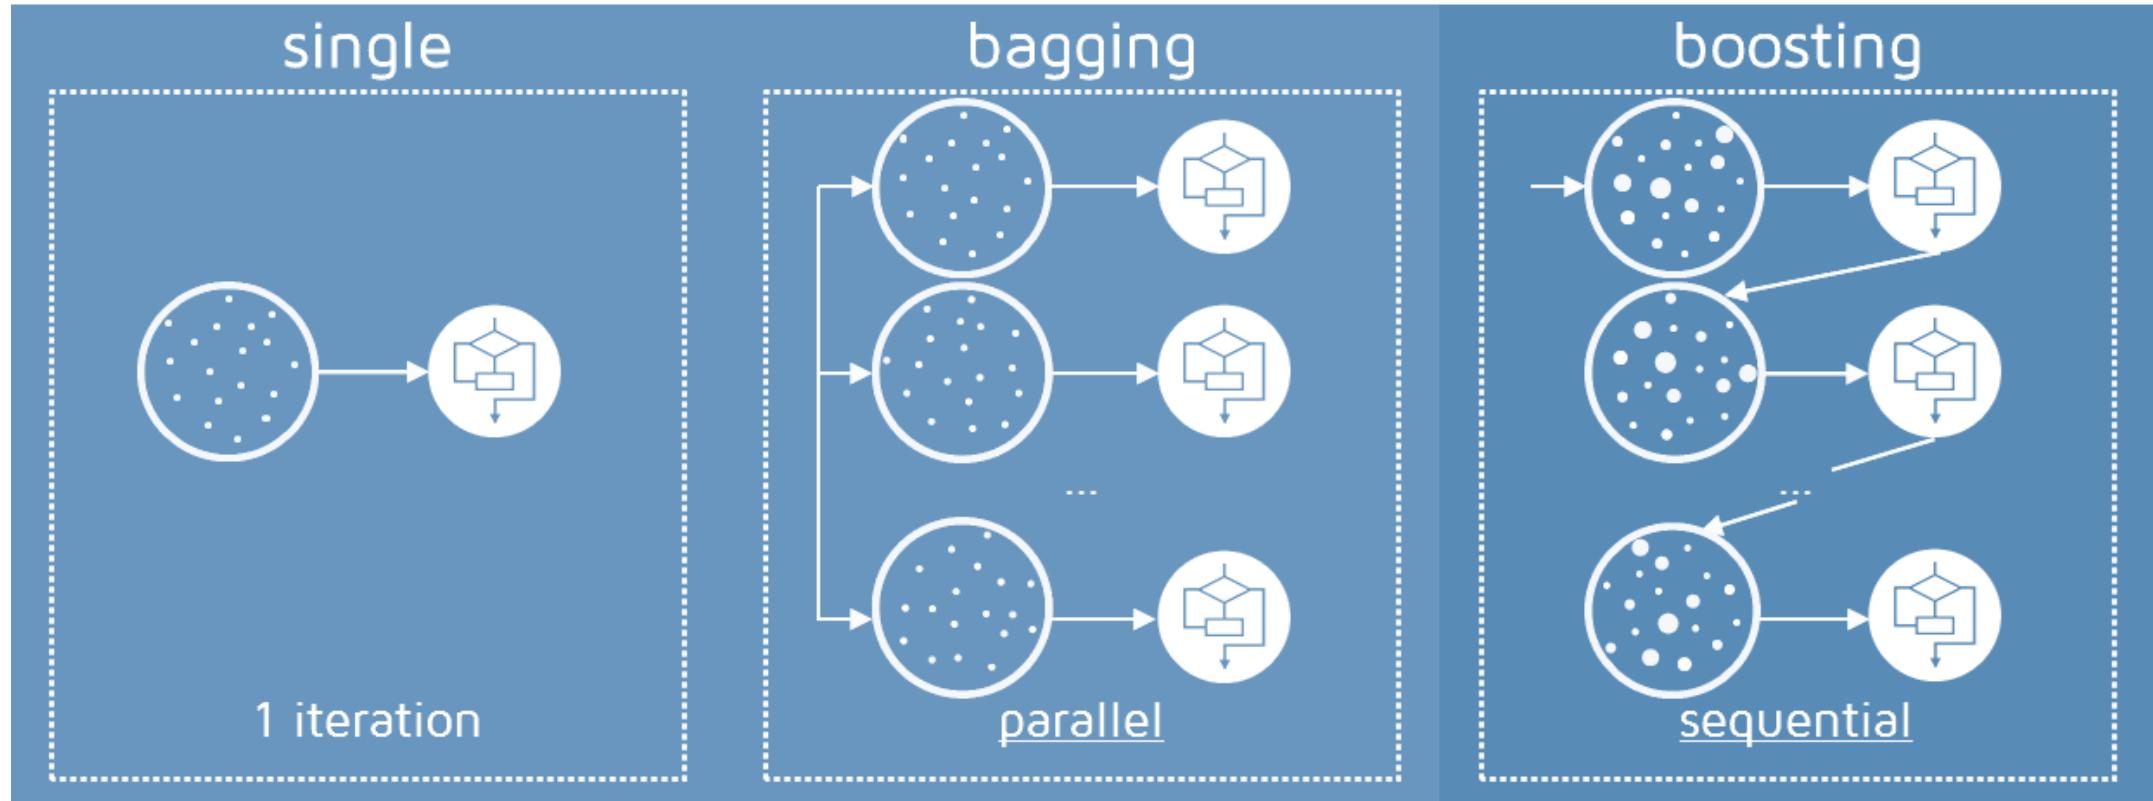
Ensemble for Generalization

- Train on random subsets of the full training data. This procedure is known as **bagging**.
- Train networks with different architectures (e.g. different numbers of layers or units, or different choice of activation function).
- Use entirely different models or learning algorithms.

Bagging



Bagging vs. Boosting



Bagging tries to solve the over-fitting problem.
Boosting tries to reduce bias.



Thank You!