

Convolutional Neural Network



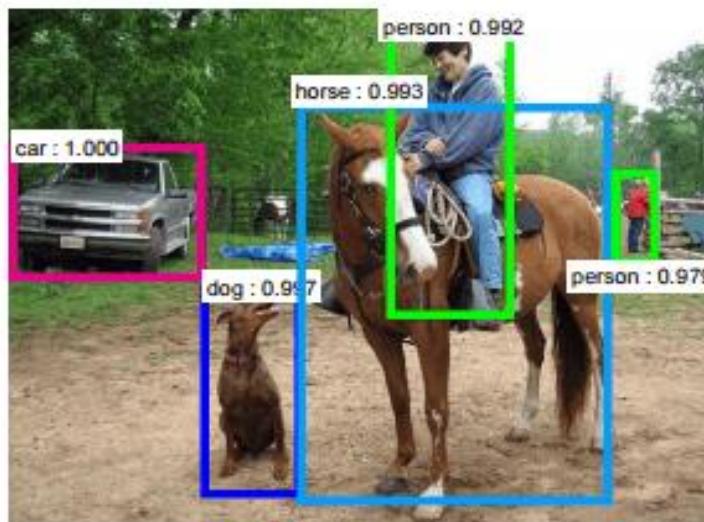
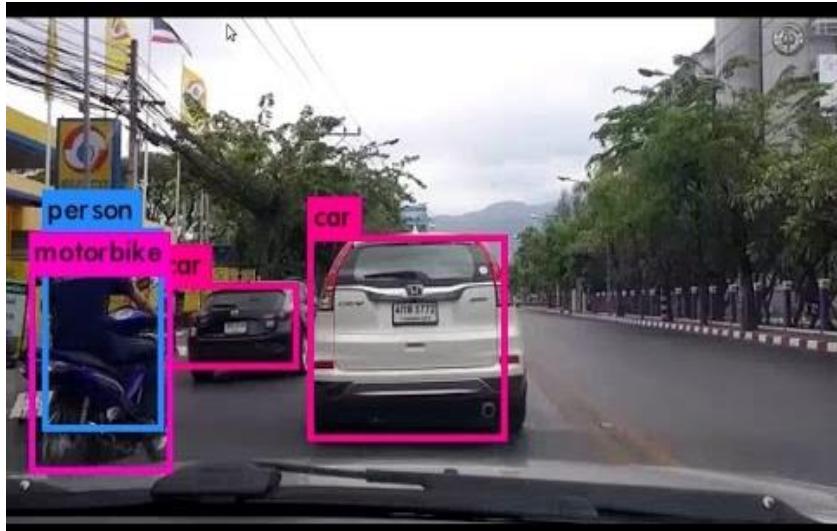
Face Recognition



Style Transferring



Image quality enhancement
Beautification

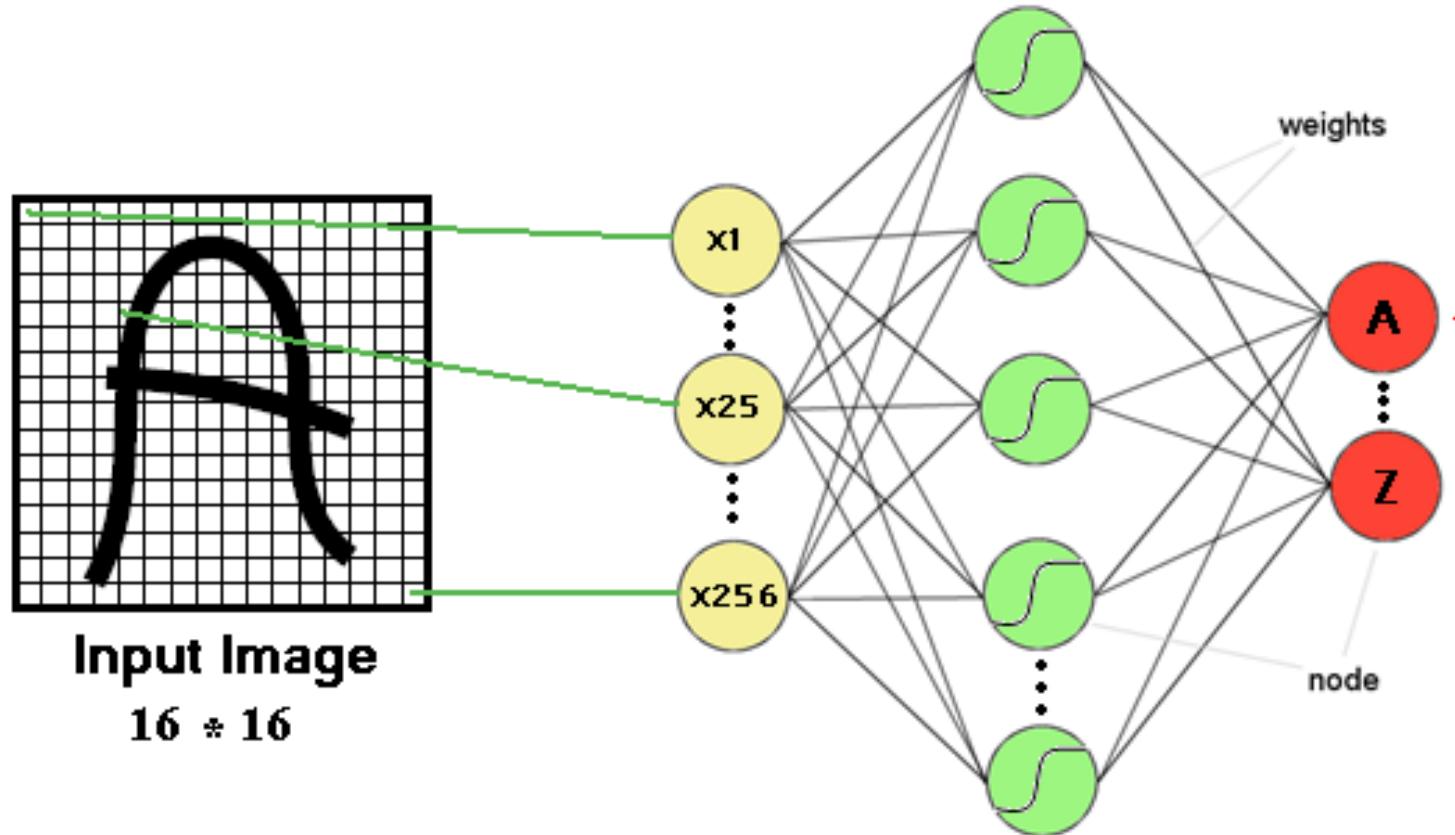


Object detection (Self driving car)

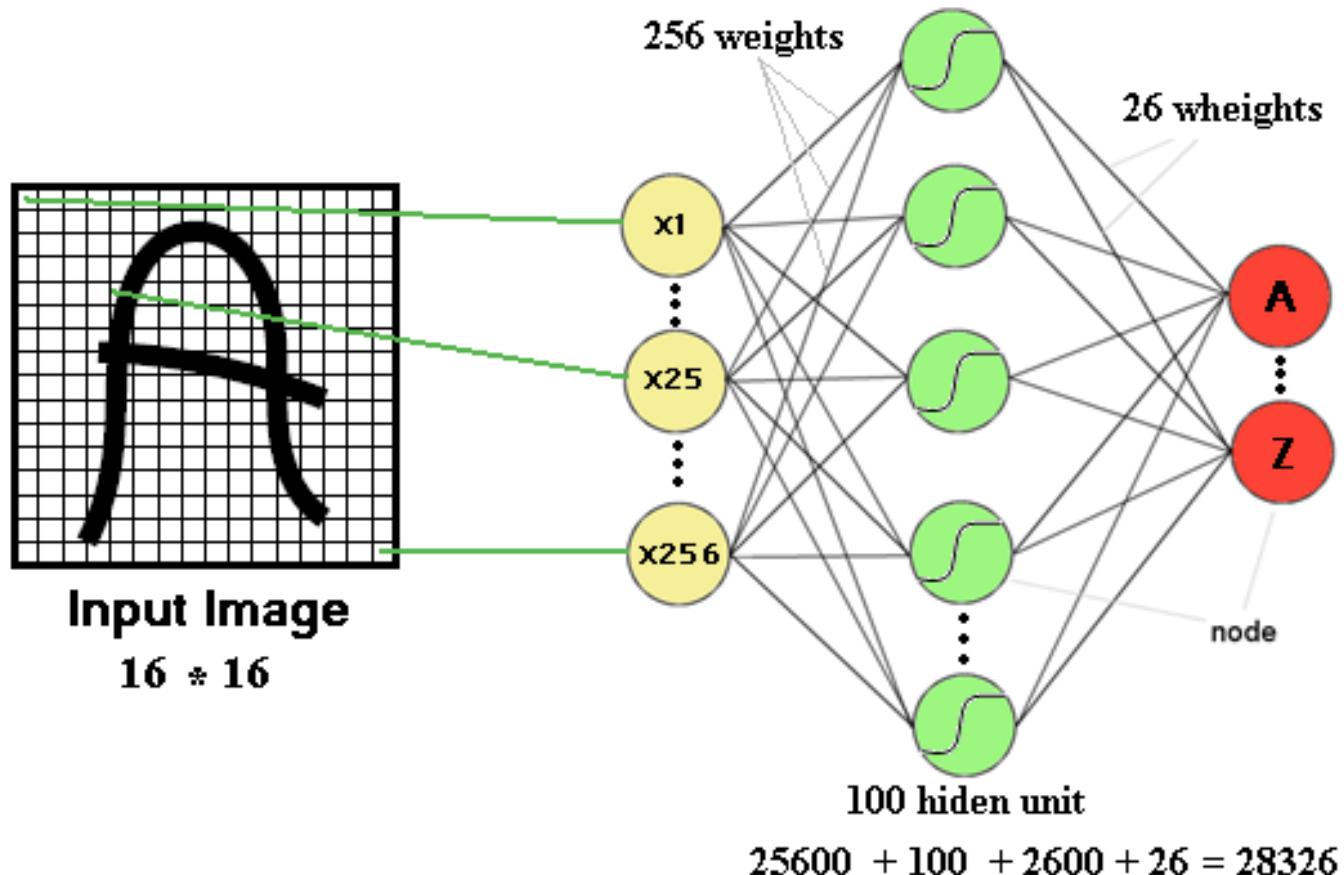
Classification

Gesture Recognition

Multi-layer perceptron and image processing



16x16 image

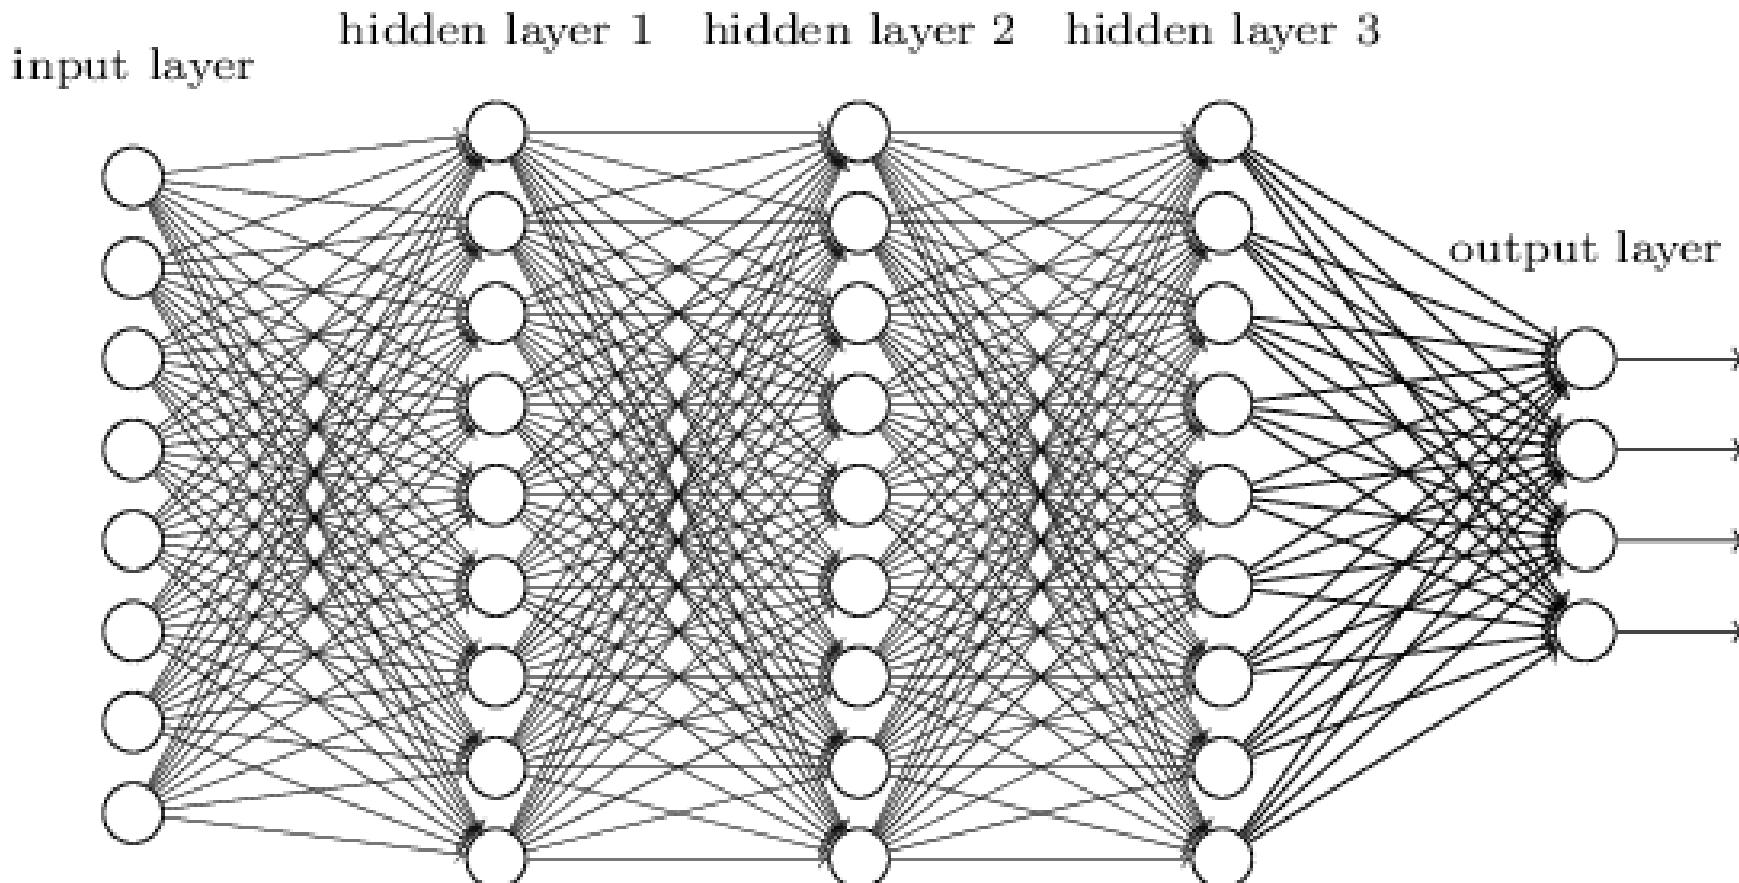


$$256 \times 100 + 100 \text{ bias} + 100 \times 26 \text{ output neurons} + 26 \text{ bias} = 28236$$

The number of **trainable parameters** becomes extremely large

ANN: Too many parameters

- We know it is good to learn a small model.
- From this fully connected model, do really need all the edges?
- Can some of these be shared?



Identify



Can we do with less information?

How much information we can throw away and still recognize the object?



10%
20%

Can we do with less information?

How much information we can throw away and still recognize the object?



10%

20%

50%

75%

Subsampling pixels will not change the object

bird



bird



Subsampling

We can subsample the pixels to make image smaller fewer parameters to characterize the image



CNNs Vs. ANNs

- ANNs suffer from **curse of dimensionality** when it comes to high resolution images
- We use filters (receptive fields) to exploit **spatial locality** by enforcing a local connectivity pattern between neurons of adjacent layers
 - *Parameter Sharing*
 - *Sparsity of connection*

Convolution

- Convolution is a pointwise multiplication of two functions to produce a third function.
- Primary purpose of convolution in CNN is to extract features from the input image.
- Matrix formed by sliding the filter over the image and computing the dot product is called the ‘Convolved Feature’ or ‘Activation Map’ or the ‘Feature Map’.

Convolution Example

3 1	0 0	1 -1	2	7	4
1 1	5 0	8 -1	9	3	1
2 1	7 0	2 -1	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

6X6 Matrix (nXn)

Convolution

*

1	0	-1
1	0	-1
1	0	-1

3X3 Filter (fXf)

=

-5	-4	0	8

(n-f+1)X(n-f+1)

Convolution Example

3	0 1	1 0	2 -1	7	4
1	5 1	8 0	9 -1	3	1
2	7 1	2 0	5 -1	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

6X6 Matrix (nXn)

Convolution

*

1	0	-1
1	0	-1
1	0	-1

3X3 Filter (fXf)

=

-5	-4	0	8

(n-f+1)X(n-f+1)

Convolution Example

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

*<http://ufldl.stanford.edu/tutorial/supervised/FeatureExtractionUsingConvolution/>

Detecting Vertical edges

$$6*6 = 36$$

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*

1	0	-1
1	0	-1
1	0	-1



$$4*4=16$$

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



- In case of ANN # parameter to train = $36*16 = 576$
- In case of CNN # parameter to train = 9

Filter Weights

1	1	1
0	0	0
-1	-1	-1

Horizontal Filter

1	0	-1
2	0	-2
1	0	-1

Sobel Filter

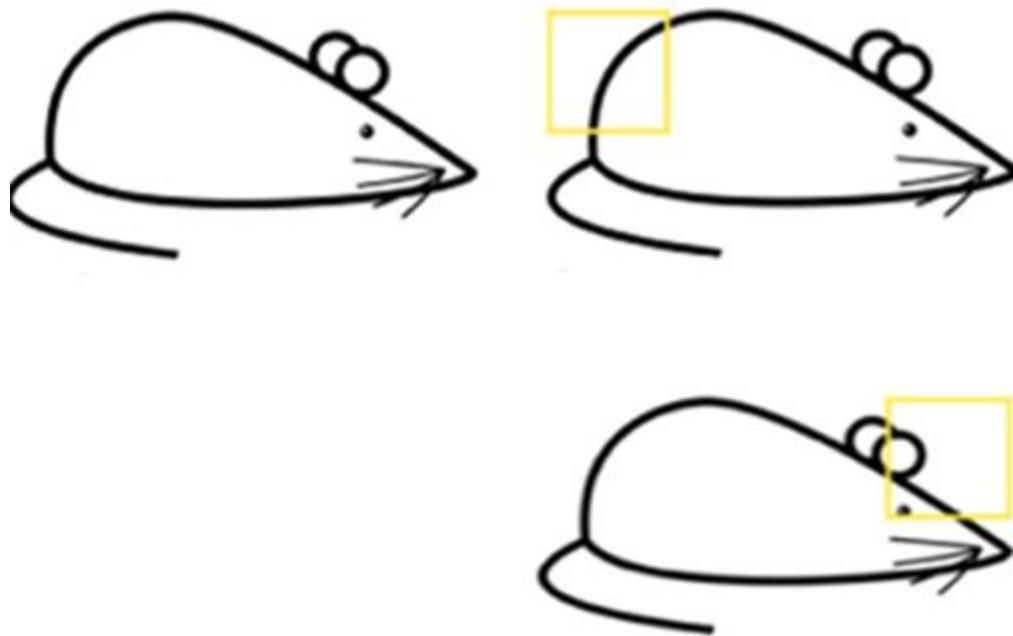
3	0	-3
10	0	-10
3	0	-3

Schorr Filter

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

Convolutional Neural Networks automatically estimates the weights of the filter

More Intuition



0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation
of a filter

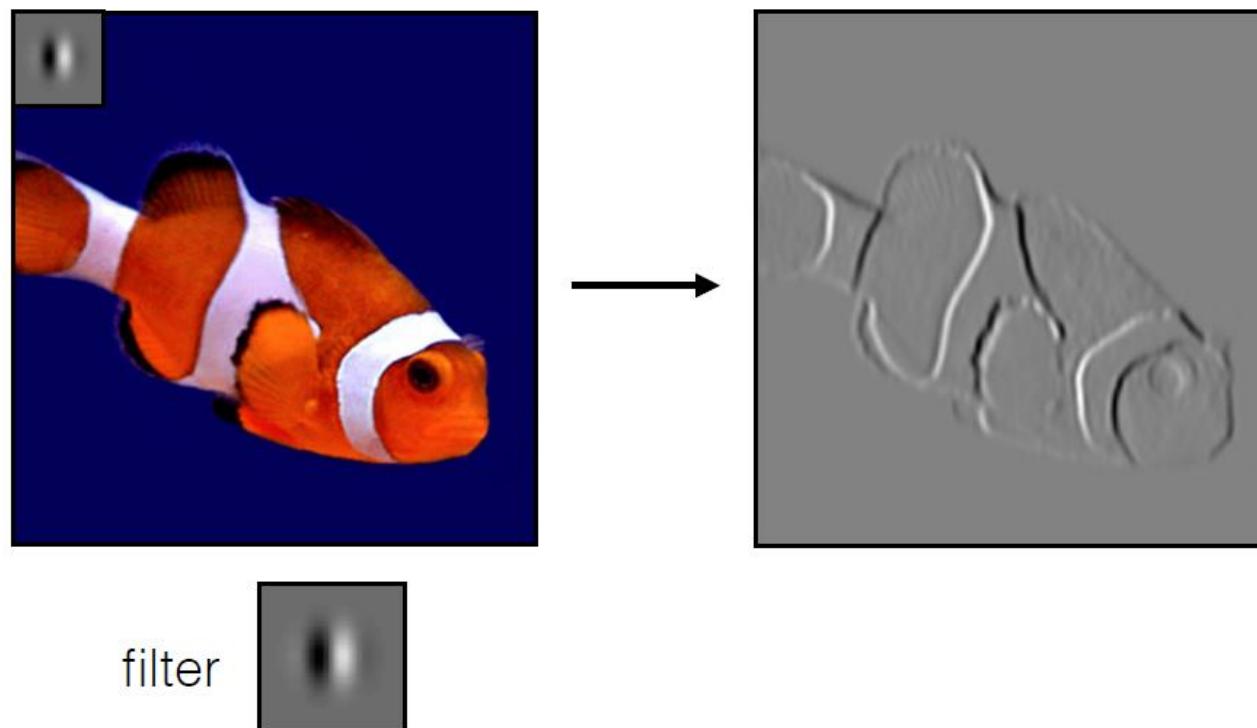


Edge

Interpretation

Convolution is just another way of computing $W^T X$

In CNN, input is **image**, kernel is **convolution filter** to be learned, **response** is the **feature map**

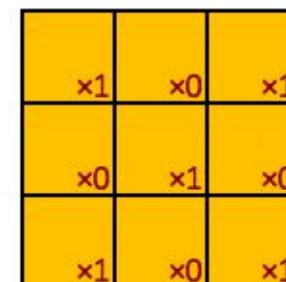


Padding

- Padding is used to preserve the original dimensions of the input
- Zeros are added to outside of the input
- Number of zero layers depend upon the size of the kernel

0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	1	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

5X5 (with padding)

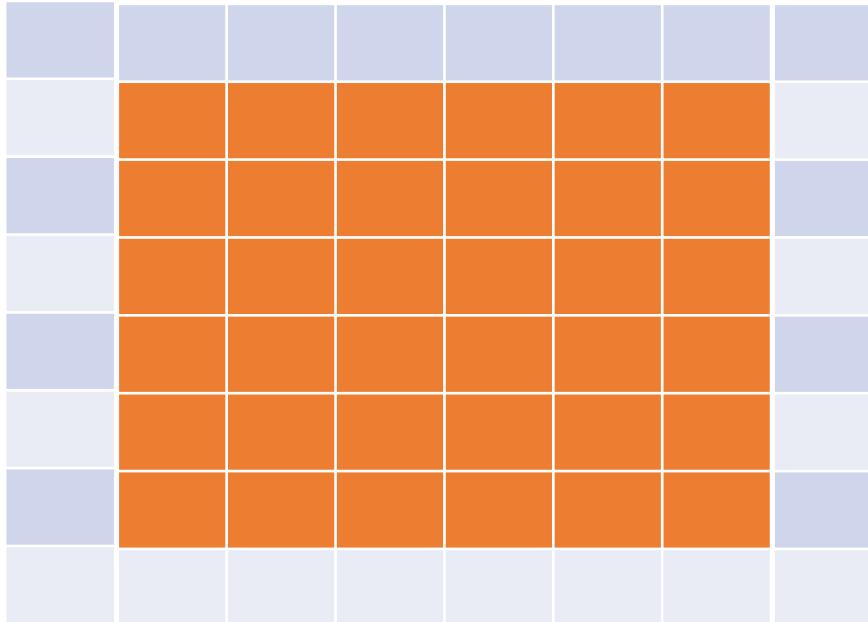


2	2	3	1	1
1	4	3	4	1
2	2	4	3	3
1	2	3	4	1
1	2	3	1	1

5X5

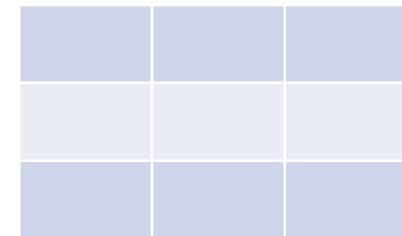
Padding

Stride=s



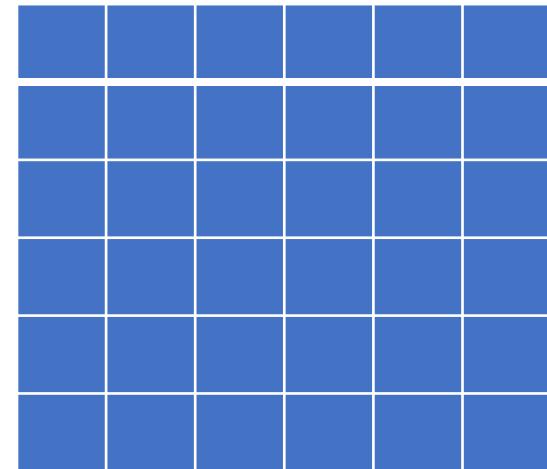
nXn 6X6 to 8X8 Padding=1

*



f 3X3

=



(n-f+1)X(n-f+1) to (n+2p-f+1)X(n+2p-f+1)
Valid to same

$$\text{Floor}\left(\frac{n+2p-f}{s} + 1\right) \times \text{Floor}\left(\frac{n+2p-f}{s} + 1\right)$$

Stride

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

6X6 Matrix

1	0	-1
1	0	-1
1	0	-1

3X3 Filter

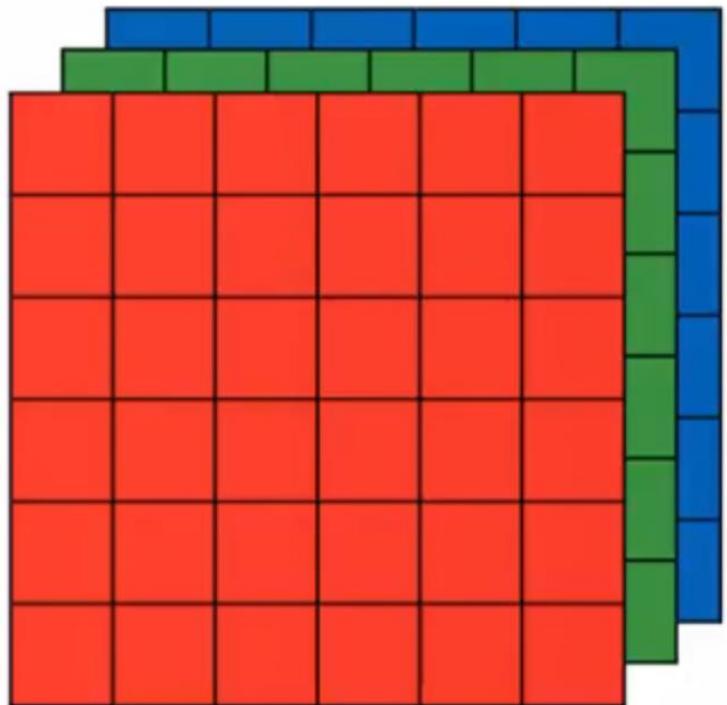
-5	8
-3	-16

2X2

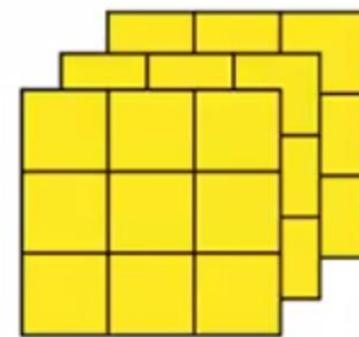
Stride=s (3 Here)

$$\text{Floor}\left(\frac{n+2p-f}{s} + 1\right) \times \text{Floor}\left(\frac{n+2p-f}{s} + 1\right)$$

Channels

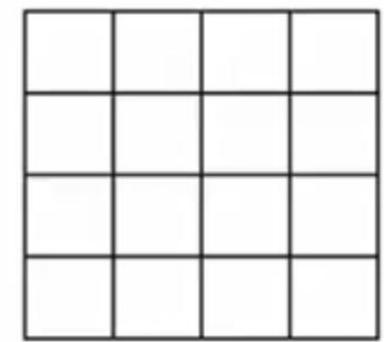


*



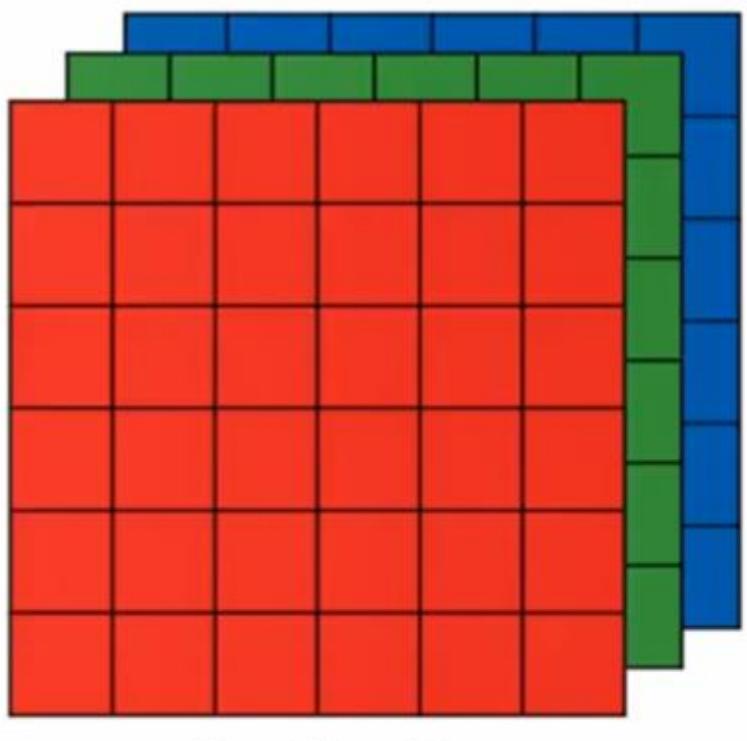
3X3 Filter

=



4 x 4

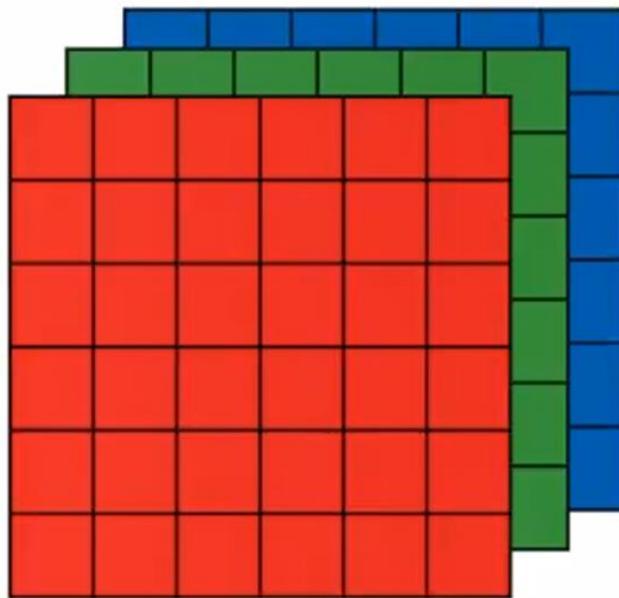
Padding =0 Stride=1



$n \times n \times n_c$

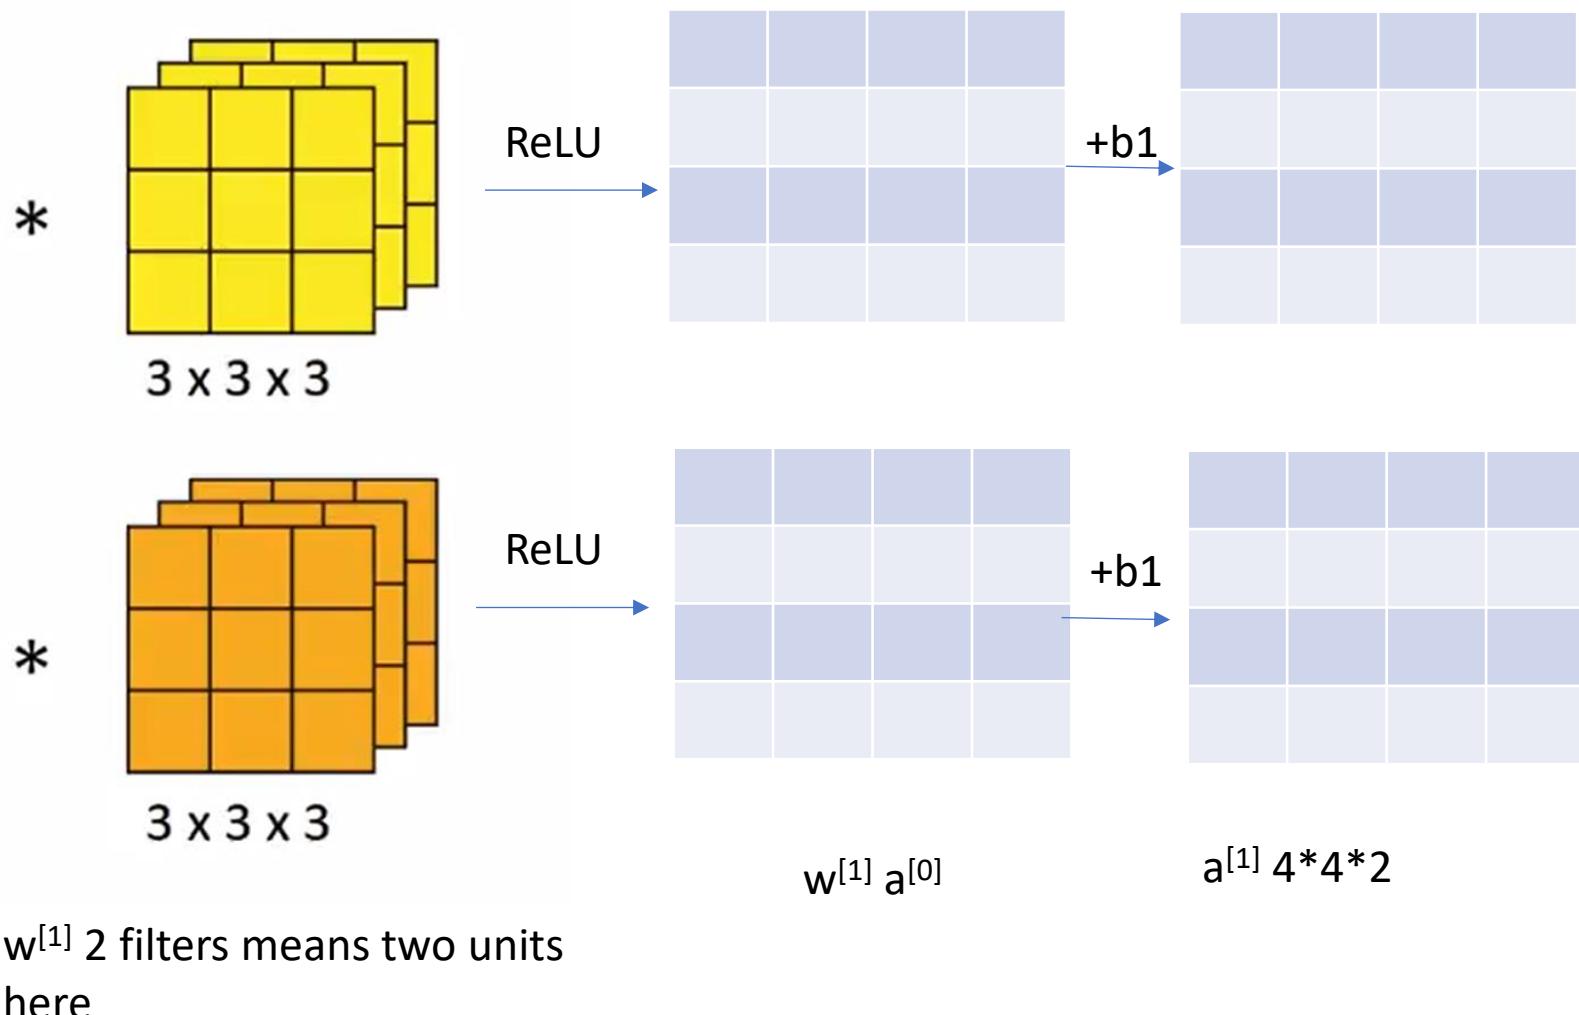
$$\begin{array}{c}
 * \\
 \text{3} \times \text{3} \times \text{3} \\
 * \\
 \text{3} \times \text{3} \times \text{3} \\
 f \times f \times n_c
 \end{array}$$

$$\begin{array}{c}
 = \\
 \text{4} \times \text{4} \\
 = \\
 \text{4} \times \text{4} \\
 n-f+1 \times n-f+1 \times n_{c'} \quad c' = \text{no of filters}
 \end{array}$$



$6 \times 6 \times 3$

$a^{[0]}$



Pooling

1	4	6	3
1	8	9	7
2	9	1	2
3	4	4	3

Max Pooling : One example of pooling layer

8	9
9	4

f=2

s=2 4X4 converted to 2X2

Function of Pooling is to progressively reduce the spatial size of the representation to reduce the amount of parameters and computation in the network.

Pooling

1	4	6	3
1	8	9	7
2	9	1	2
3	4	4	3

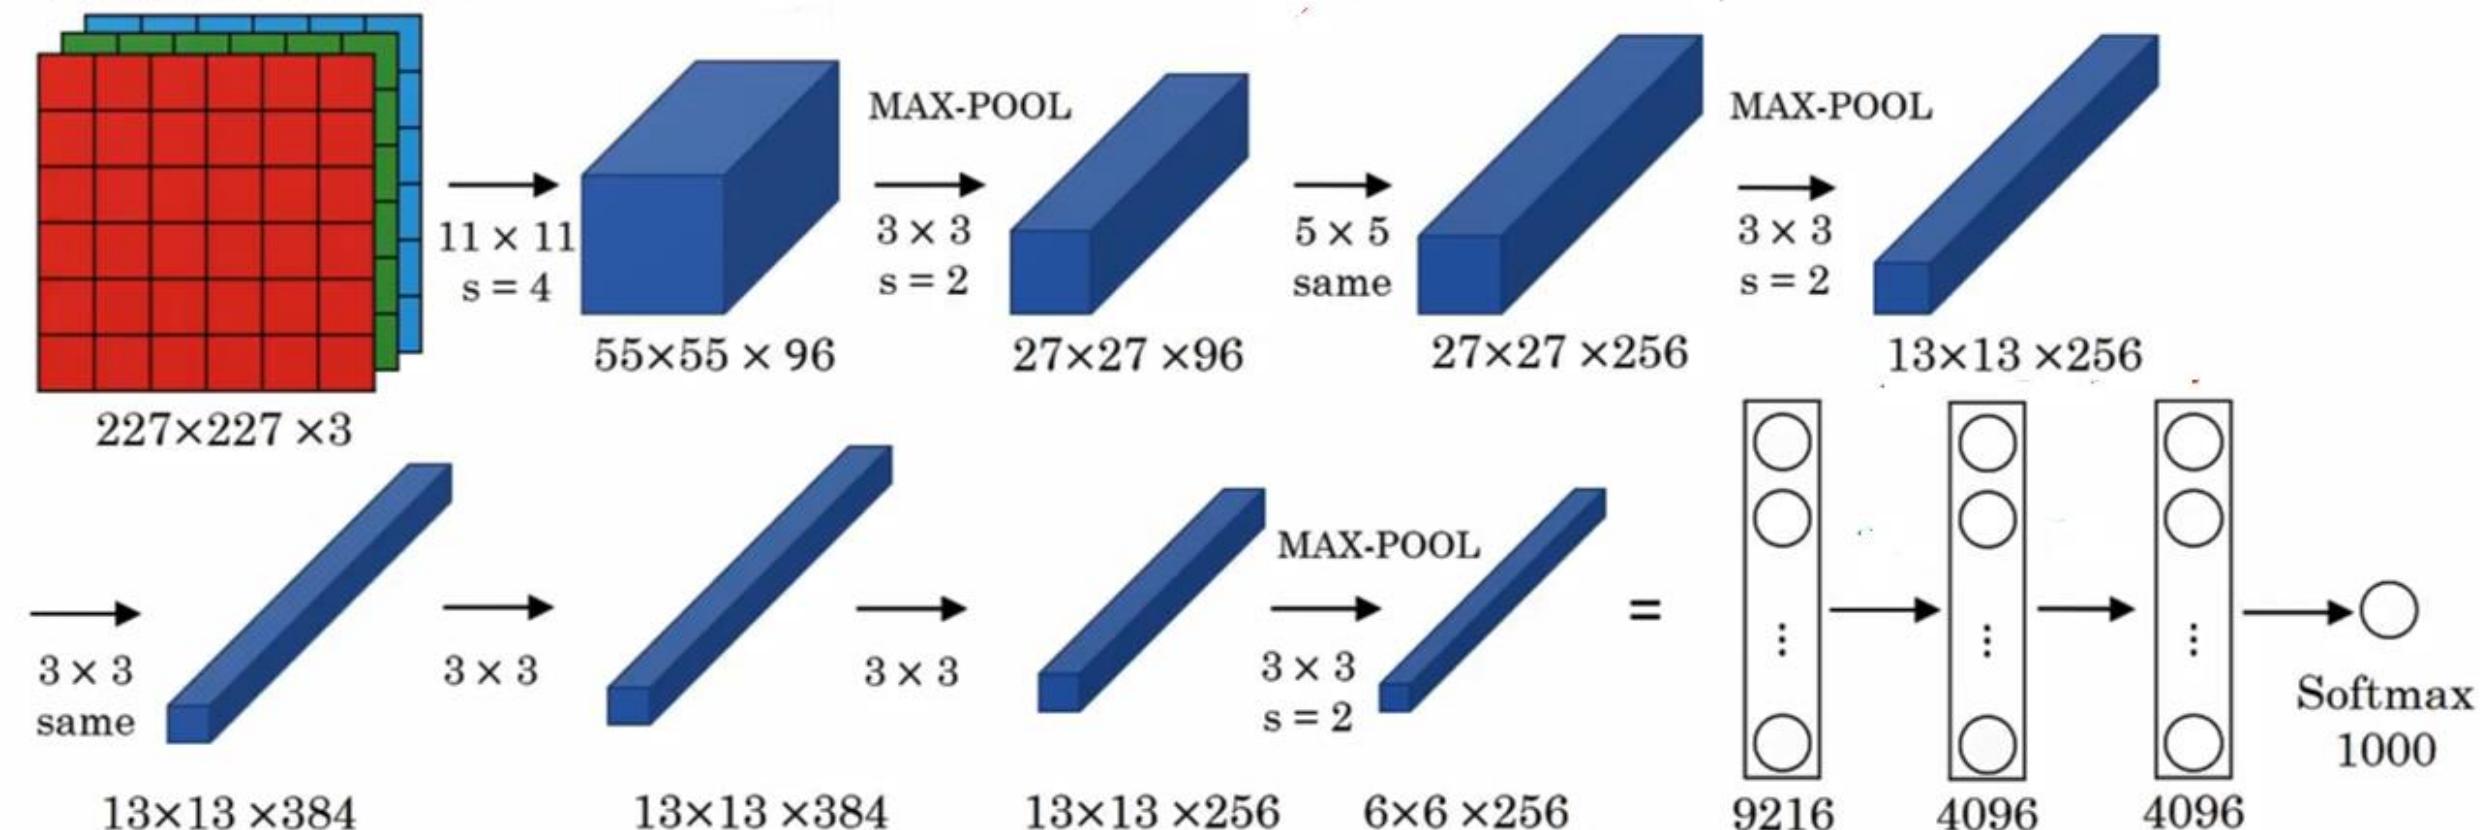
3.5	6.25
4.5	2.5

Average Pooling : Another example of pooling layer

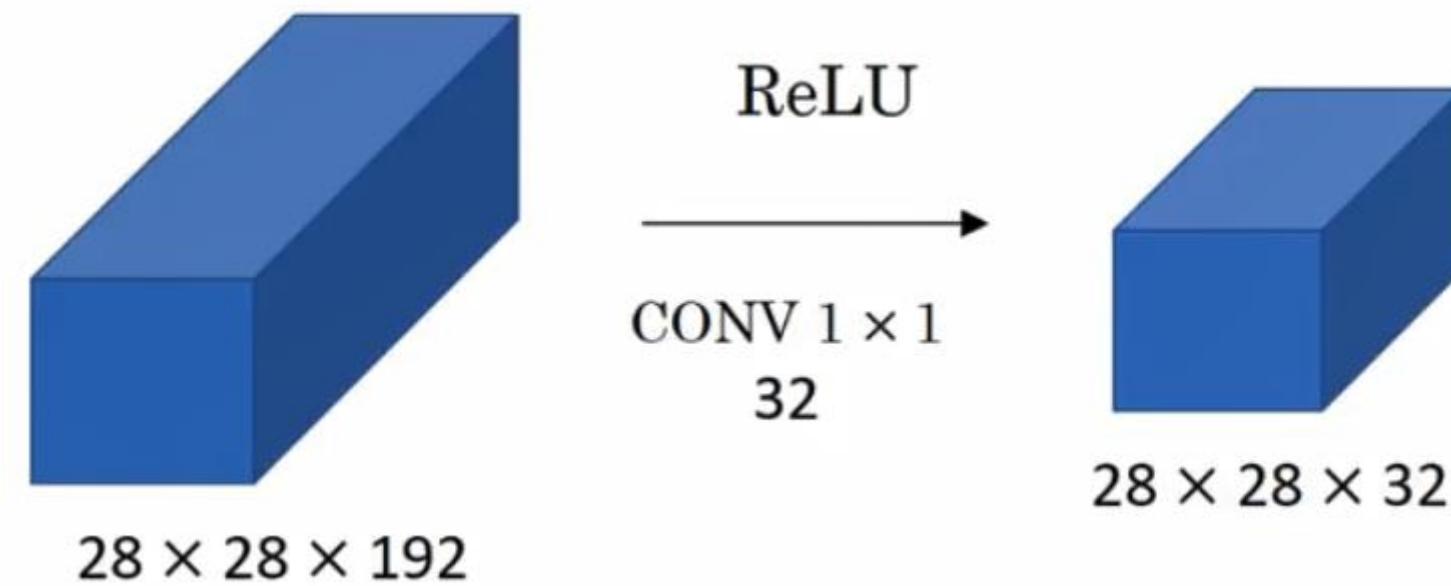
f=2

s=2 4X4 converted to 2X2

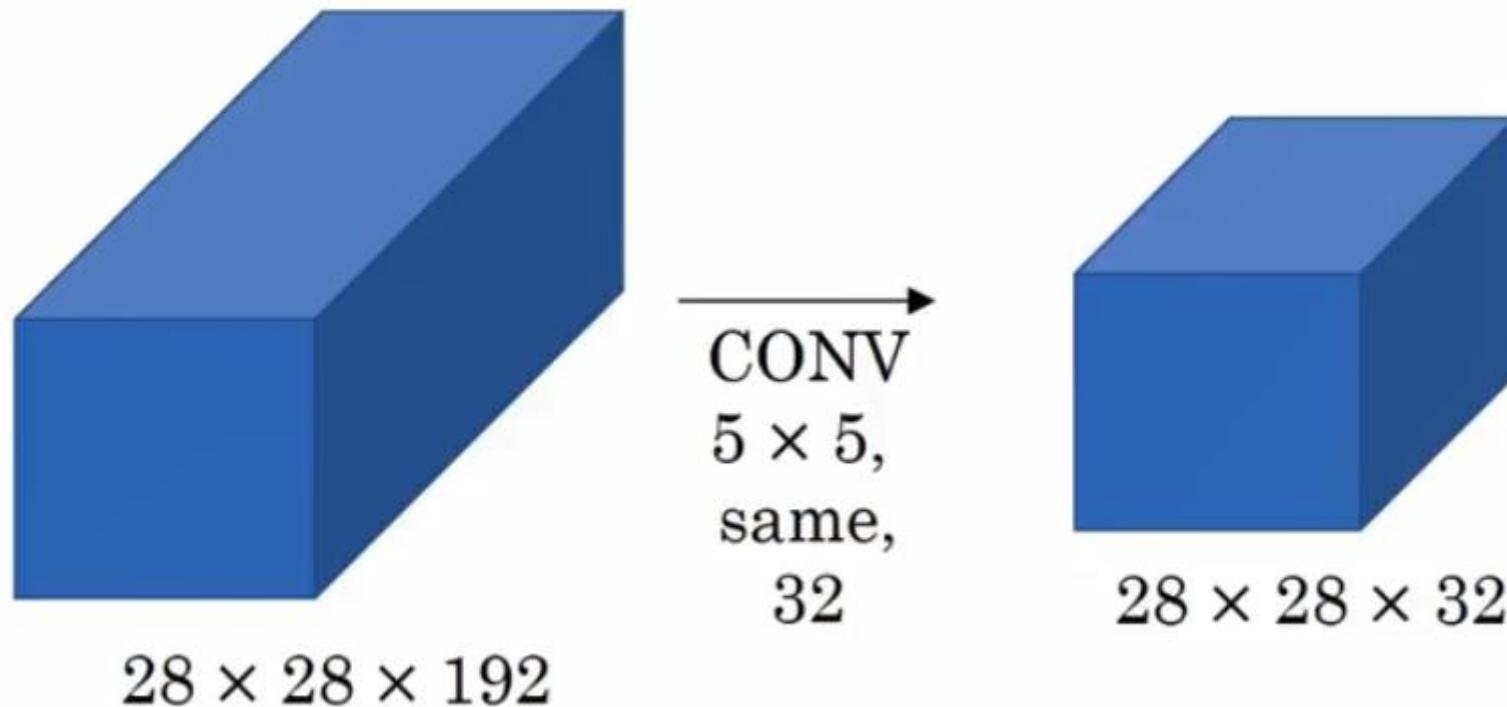
AlexNet



Shrinking of Channels

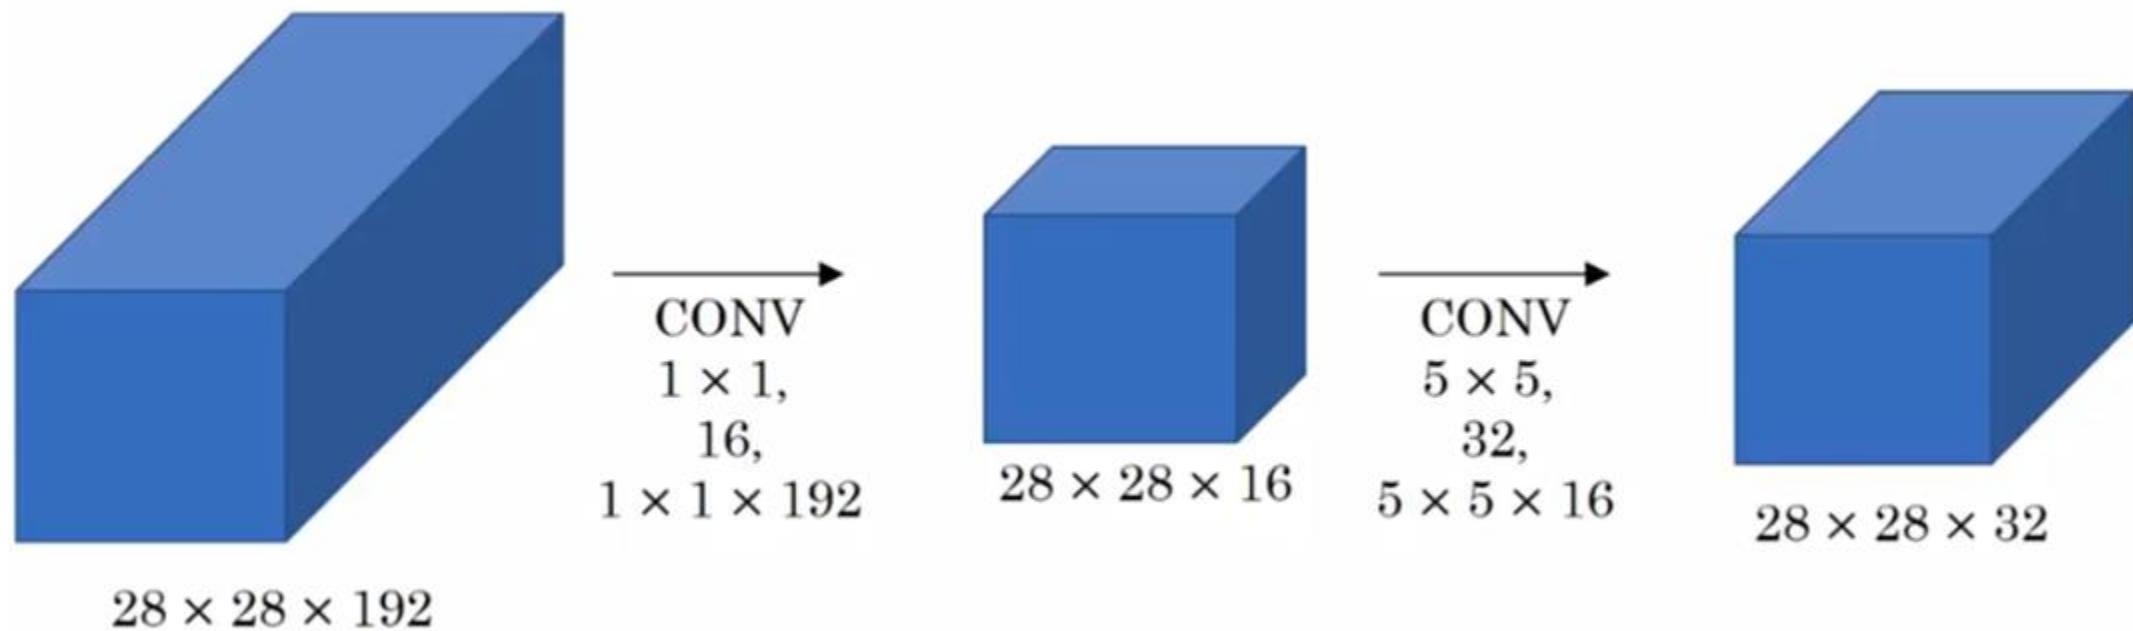


Problem of Computational Cost



$28 \times 28 \times 32 \times 5 \times 5 \times 192 = 120 \text{ Million Calculations}$

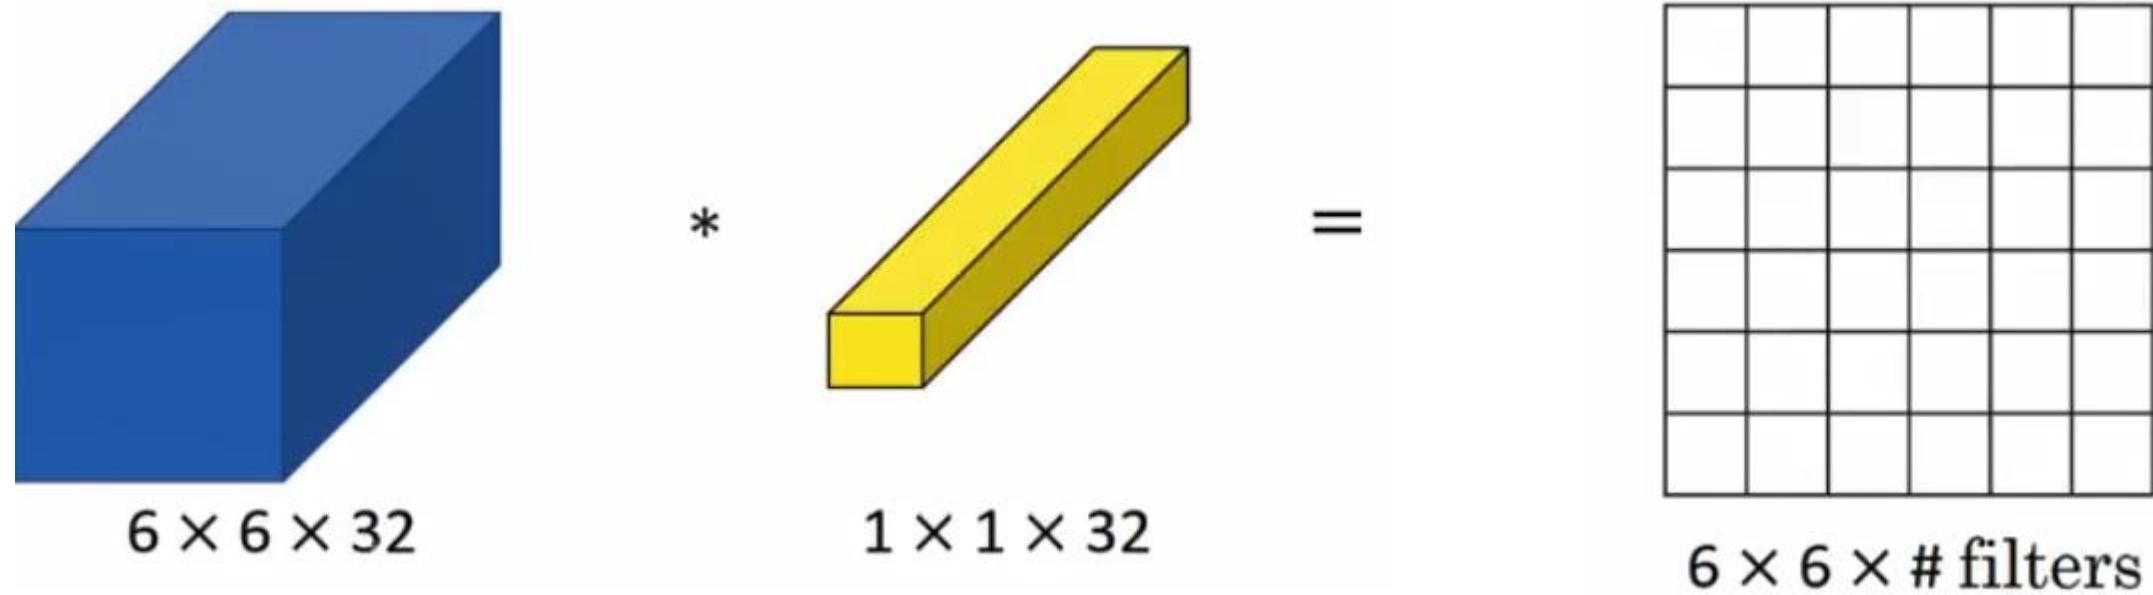
Using 1X1 Convolution to reduce the cost

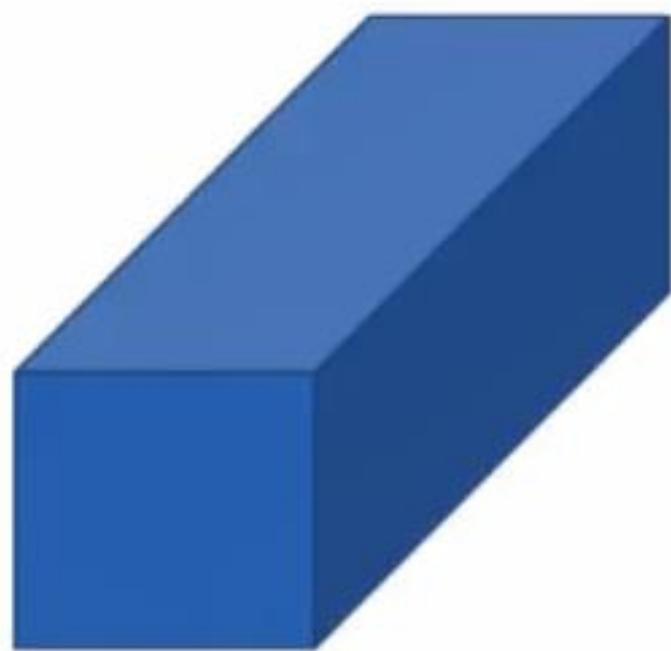


$$28 \times 28 \times 16 \times 192 + 28 \times 28 \times 32 \times 5 \times 5 \times 16 = 12.4 \text{ Million Calculations}$$

Well-known Network Architectures

Innovative use of 1X1 Convolution Filter





ReLU



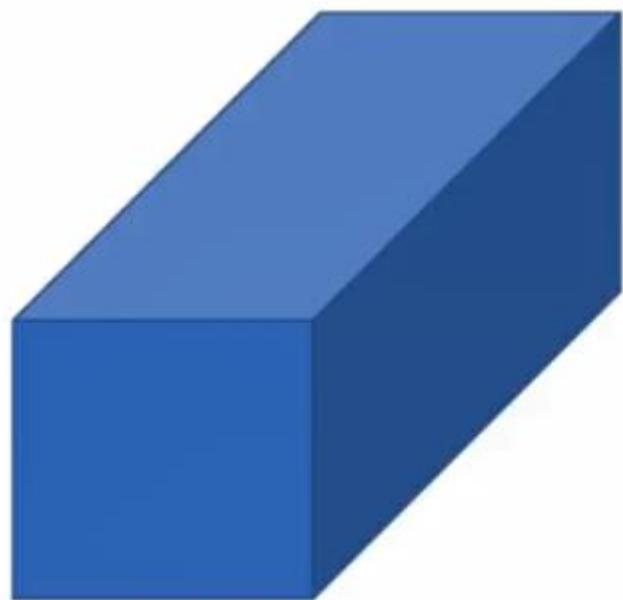
CONV 1×1
32



$28 \times 28 \times 32$

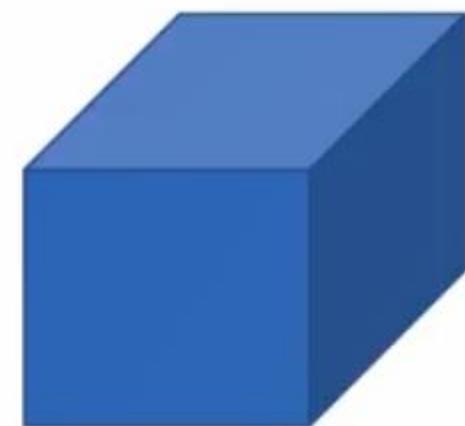
$28 \times 28 \times 192$

Fire Layer



$28 \times 28 \times 192$

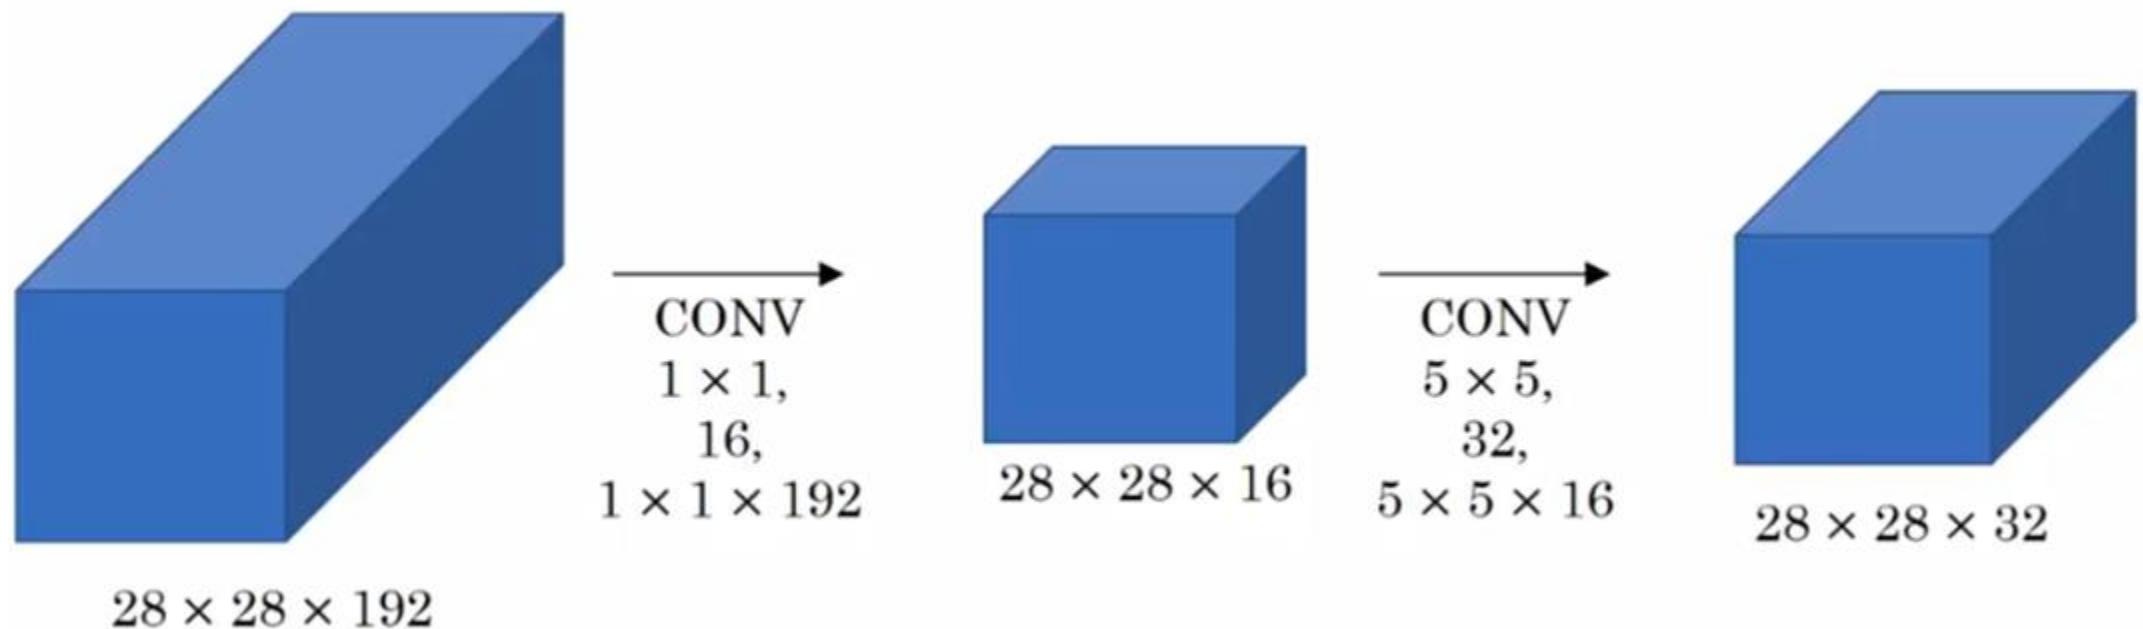
$\xrightarrow{\text{CONV}}$
 5×5 ,
same,
32



$28 \times 28 \times 32$

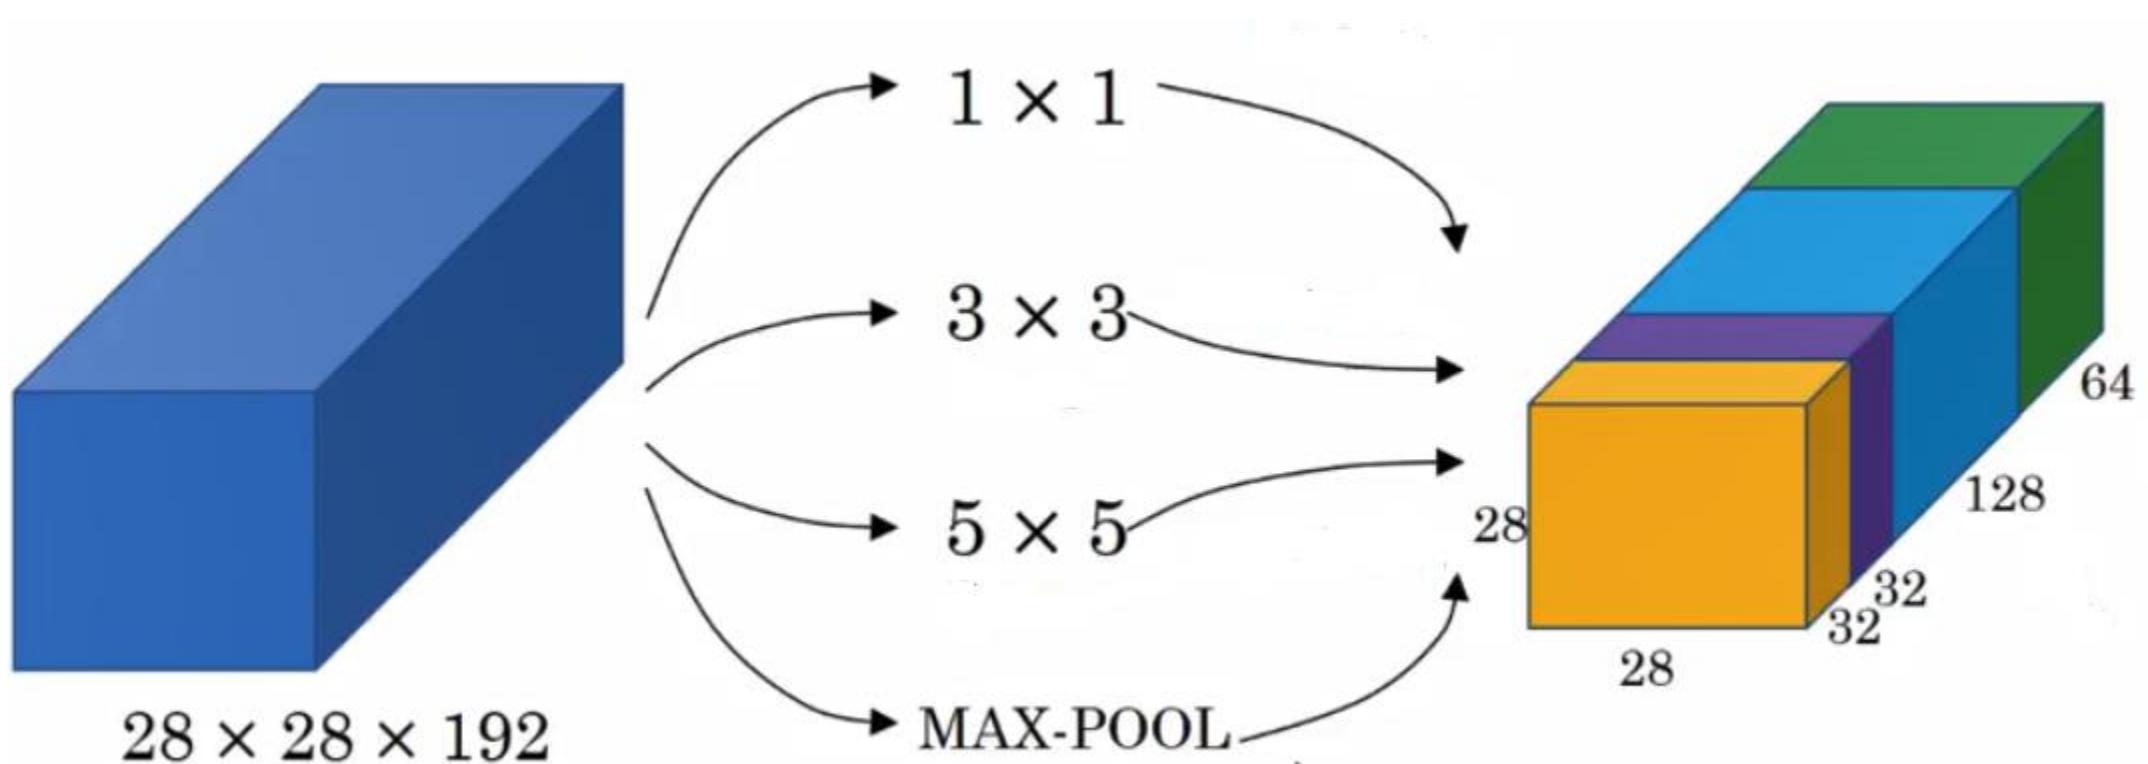
$28 \times 28 \times 32 \times 5 \times 5 \times 192 = 120$ Million Calculations
 $32 \times 5 \times 5 \times 192 = 153600$ parameters

Fire Layer

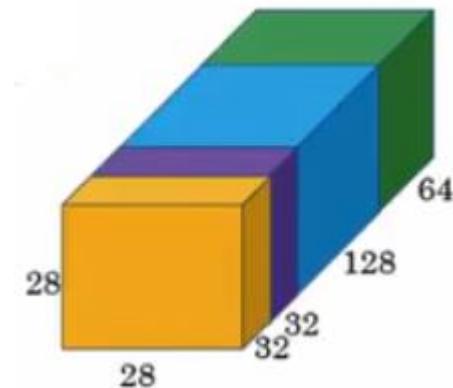
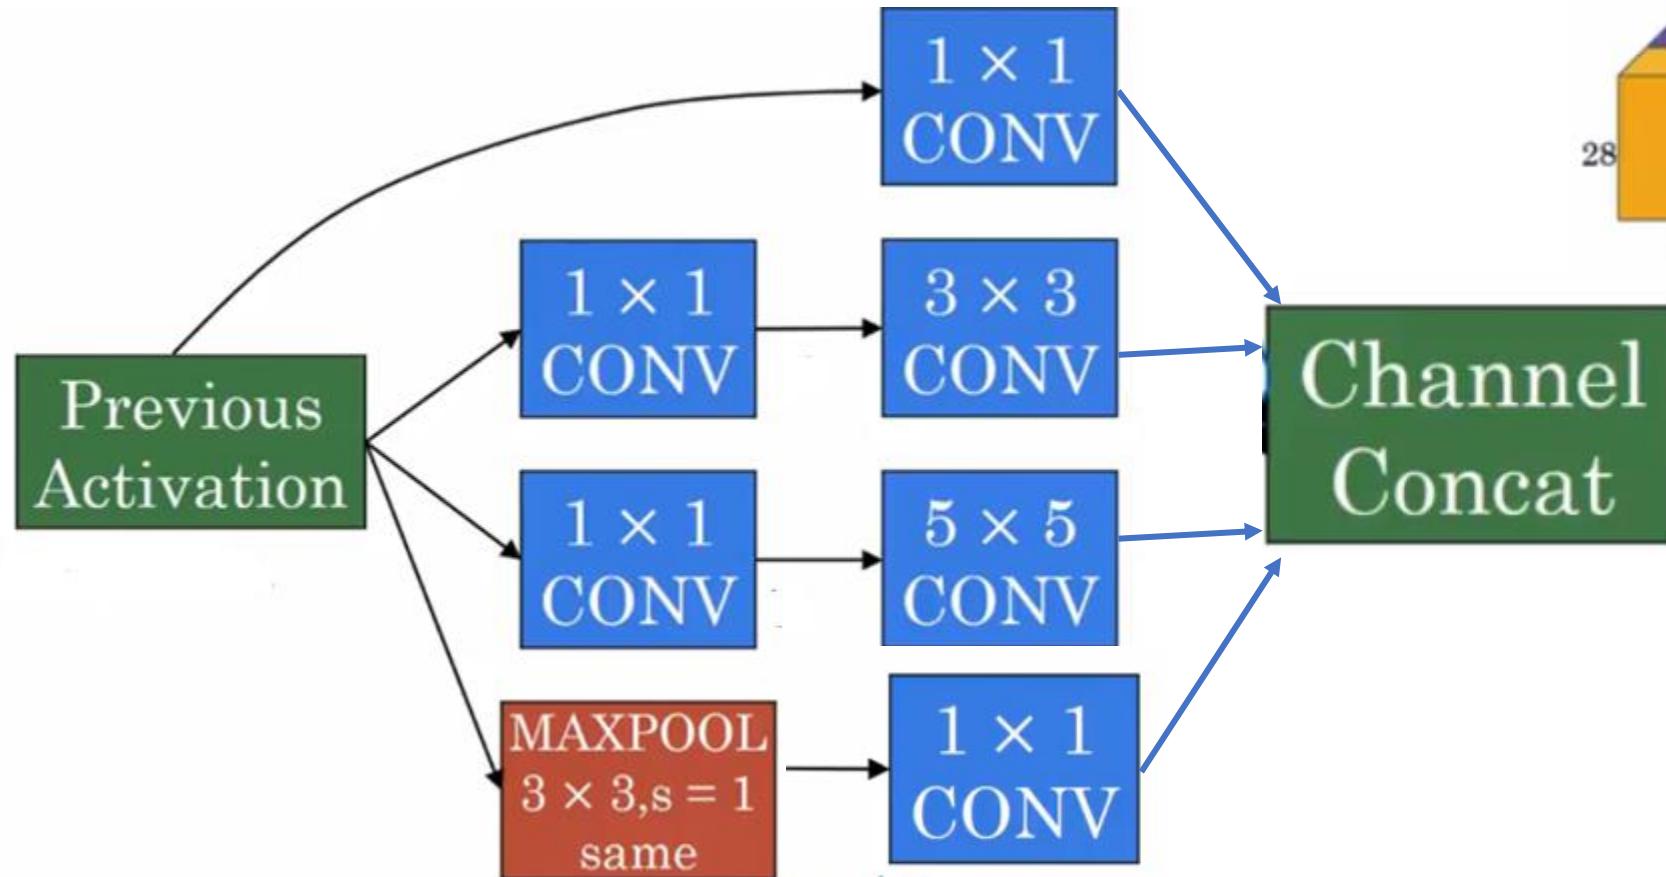


$28 \times 28 \times 16 \times 192 + 28 \times 28 \times 32 \times 5 \times 5 \times 16 = 12.4$ Million Calculations
 $16 \times 192 + 32 \times 5 \times 5 \times 16 = 15873$ parameters 10x less parameters

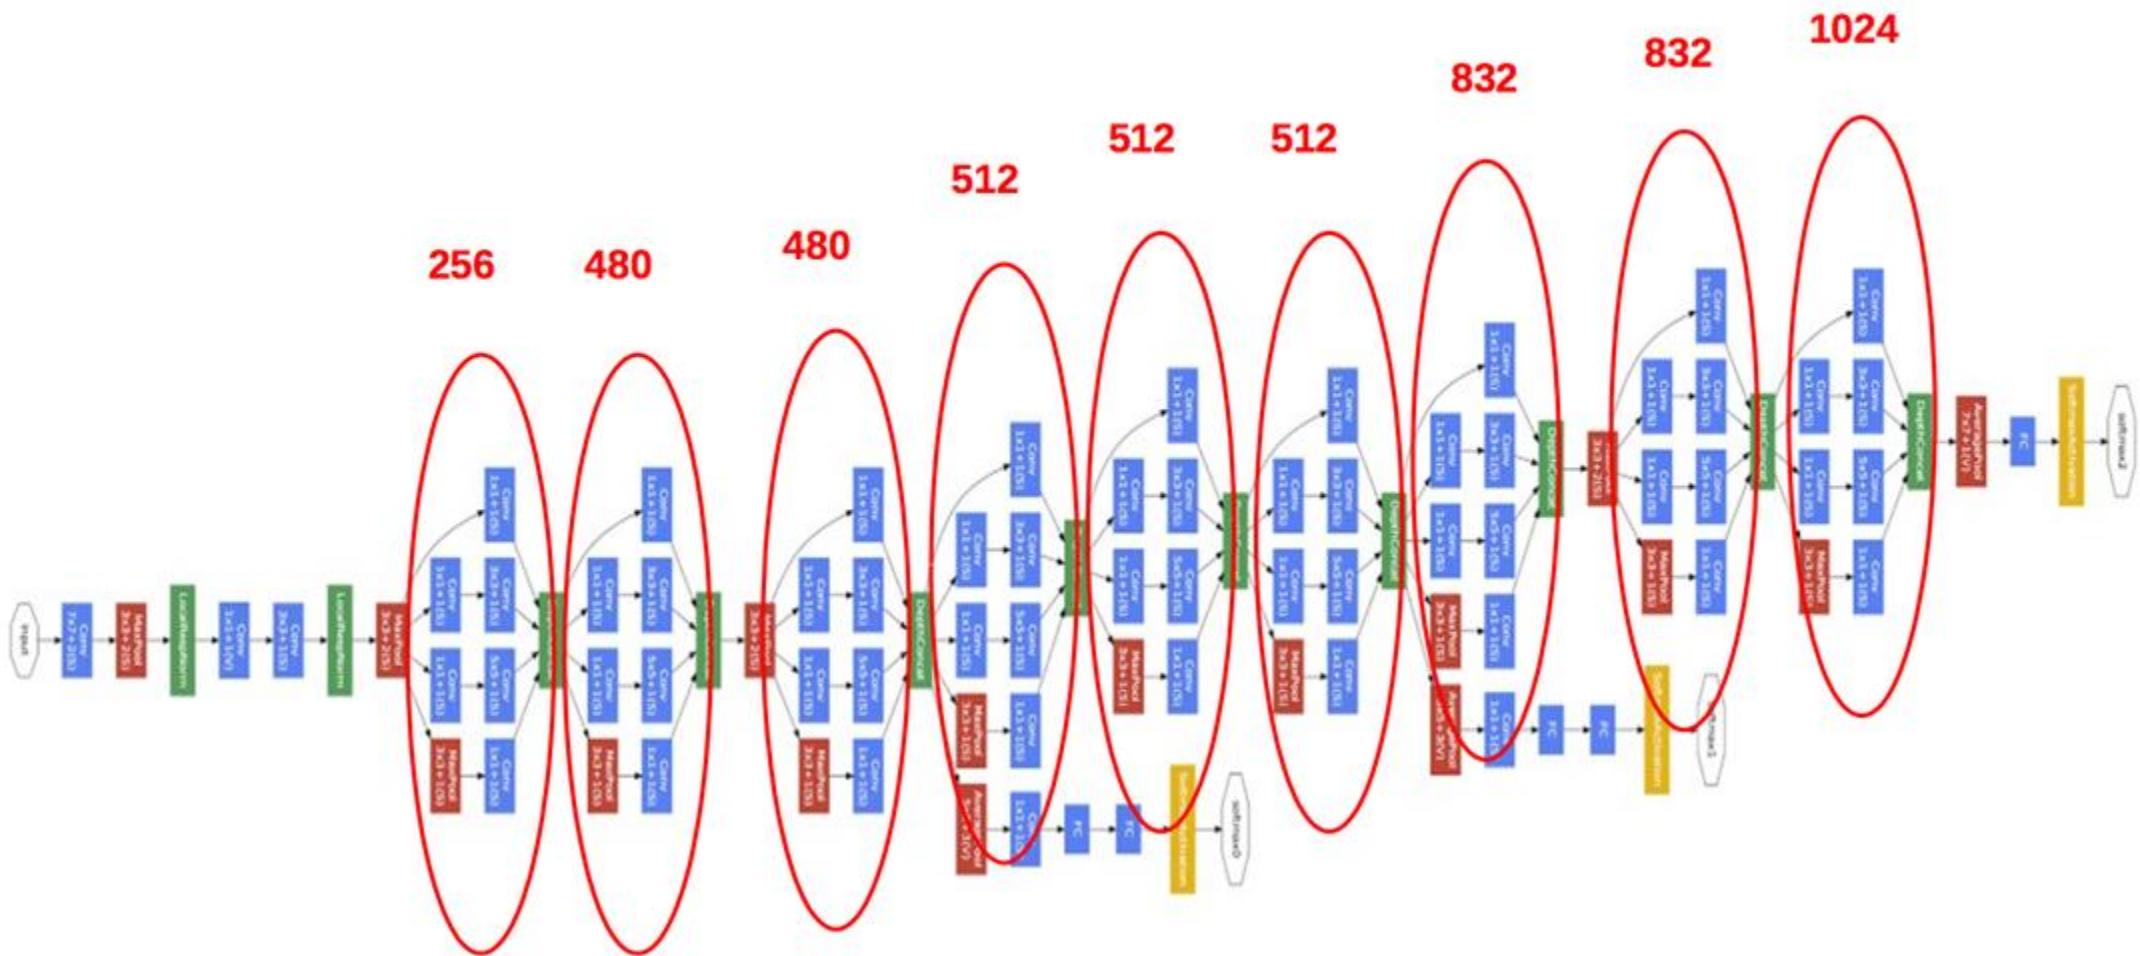
Inception Network



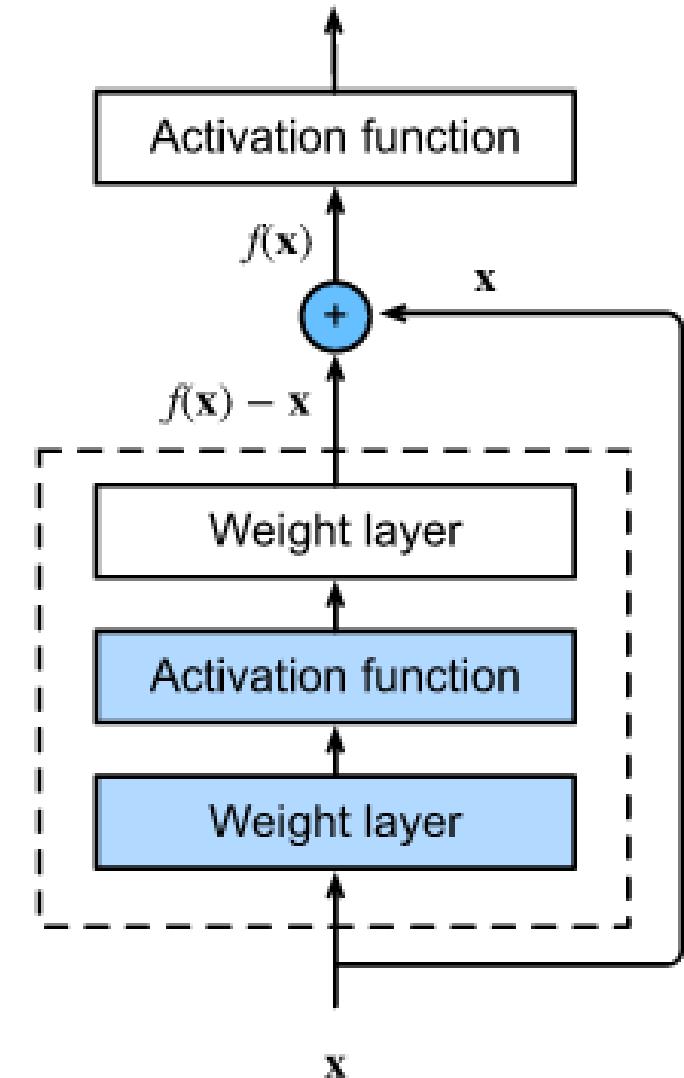
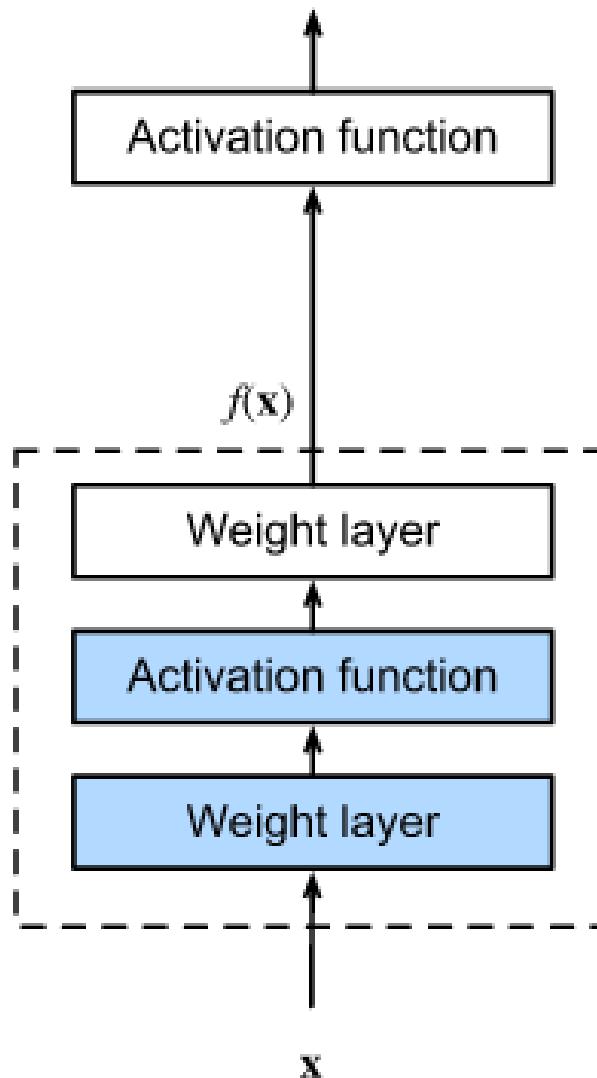
Single Unit of Inception Network



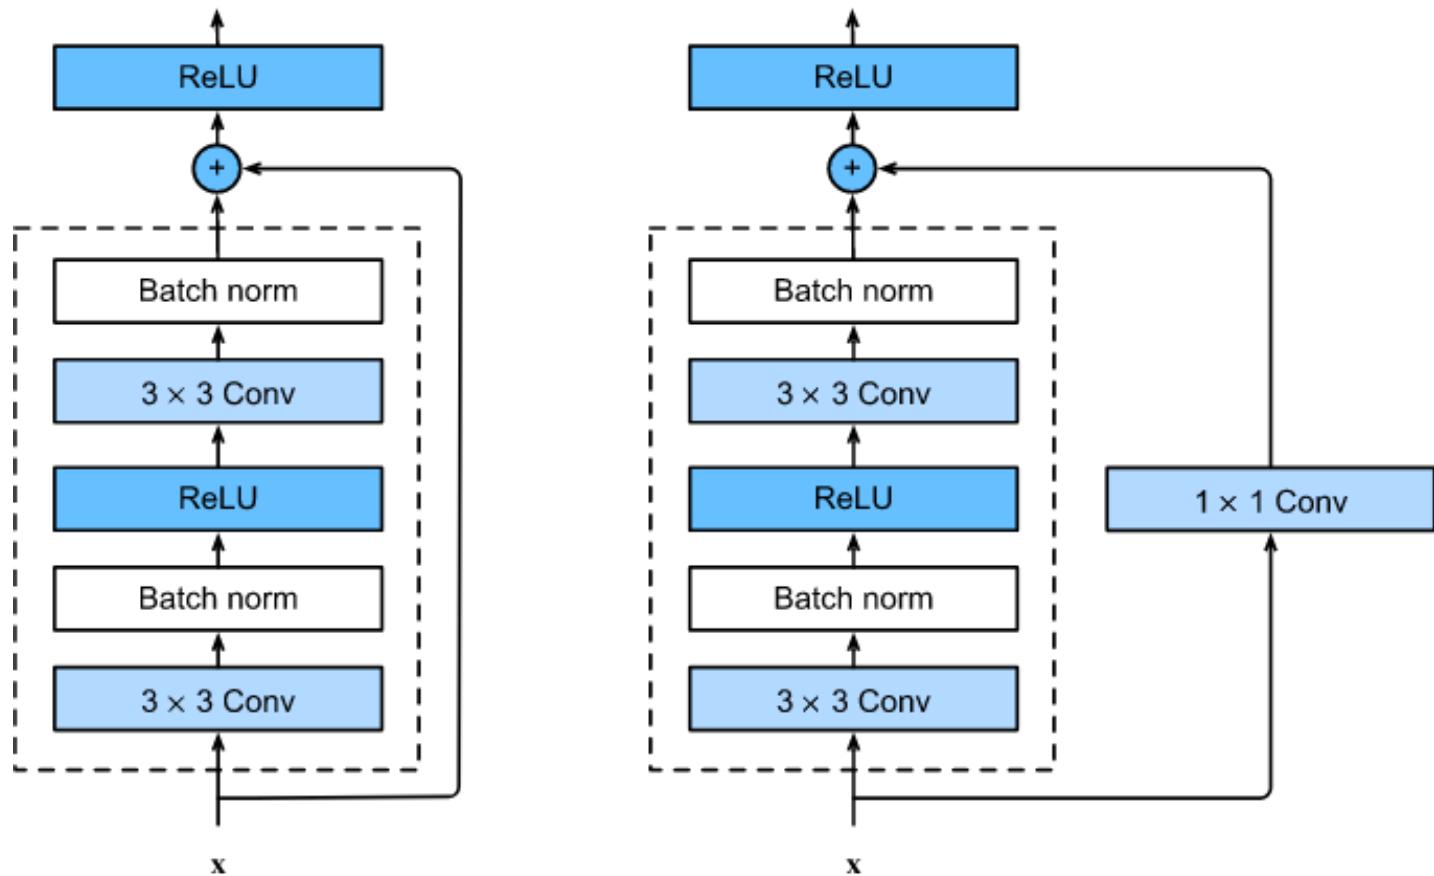
Inception Network and Inception Block



Residual Block



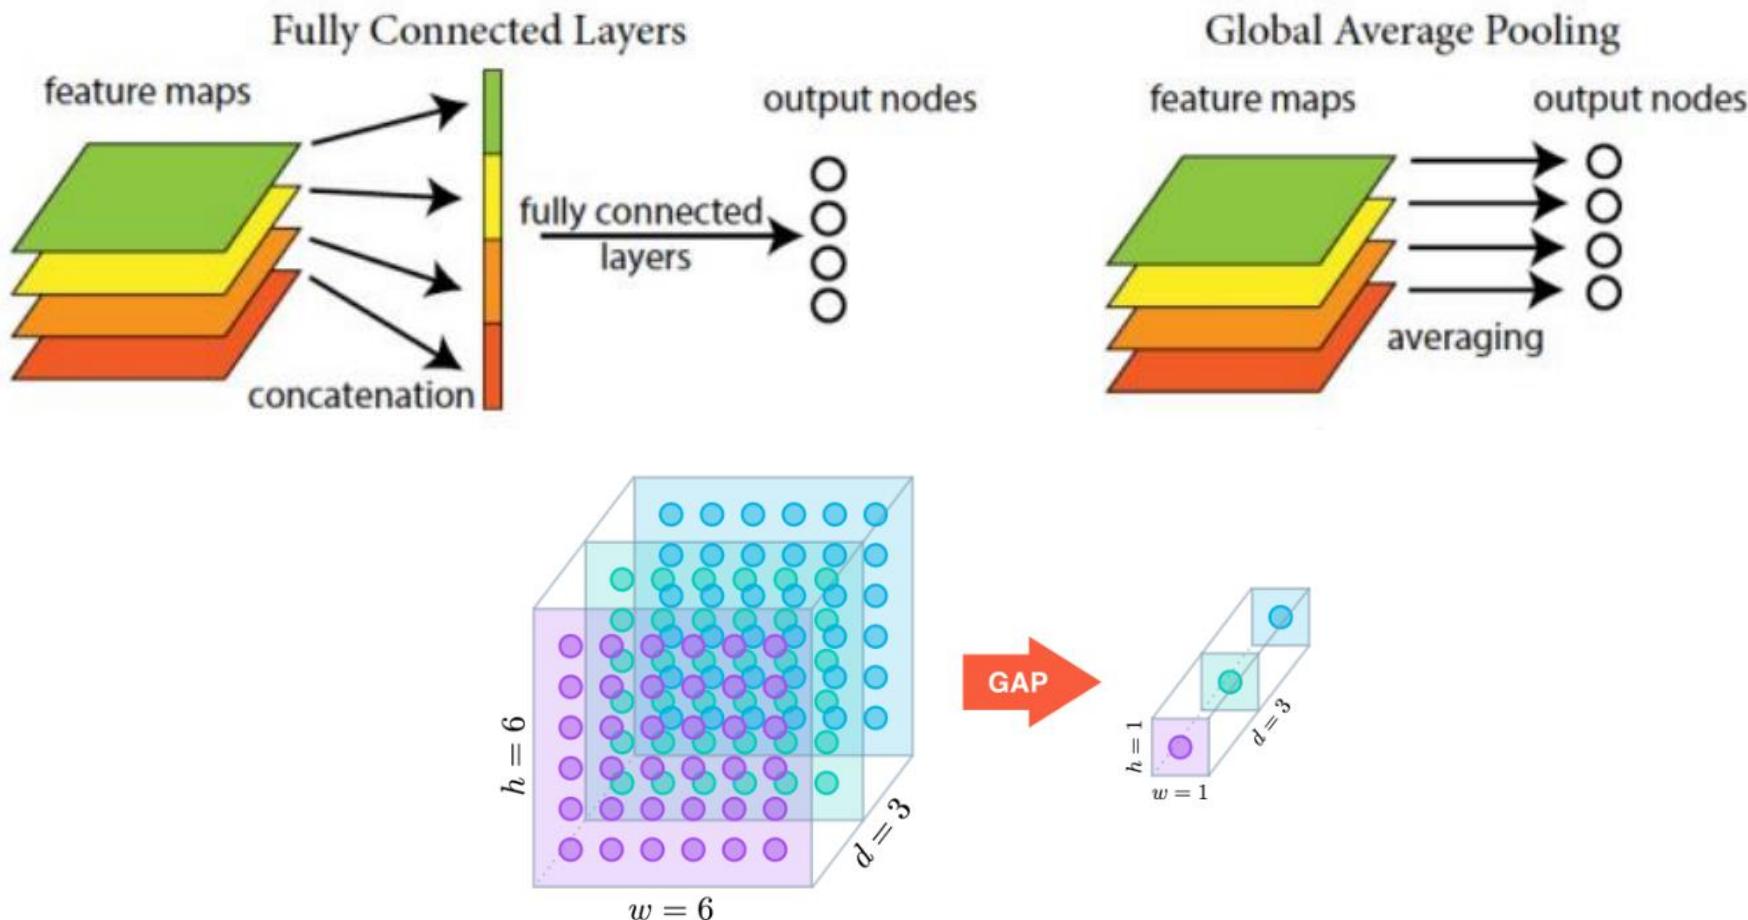
Residual Block



GoogLe Net: Global Average Pooling

- **Problems with fully connected (FC) layers:**
- More than 90% parameters of Alexnet and VGG are in the Fully Connected layers.
- One single particular layer in VGG contains 100 million parameters alone.
- Prone to overfitting.
- Heavily dependent on regularization methods like dropout.

GoogLe Net: Global Average Pooling



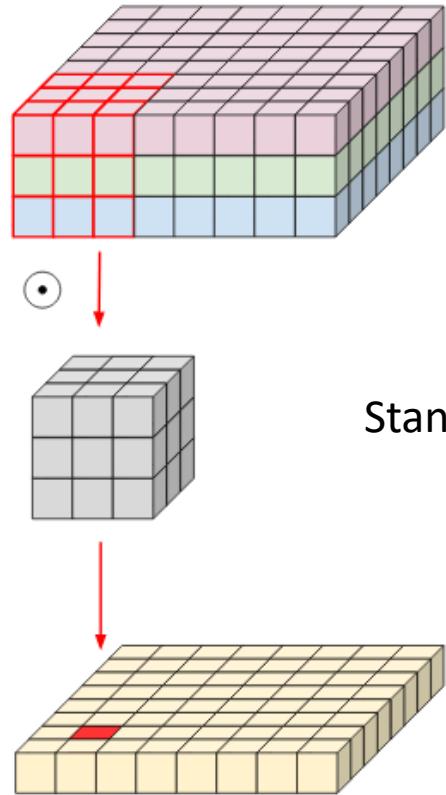
Source: <https://alexisbcook.github.io/2017/global-average-pooling-layers-for-object-localization/>

GoogLe Net: Global Average Pooling

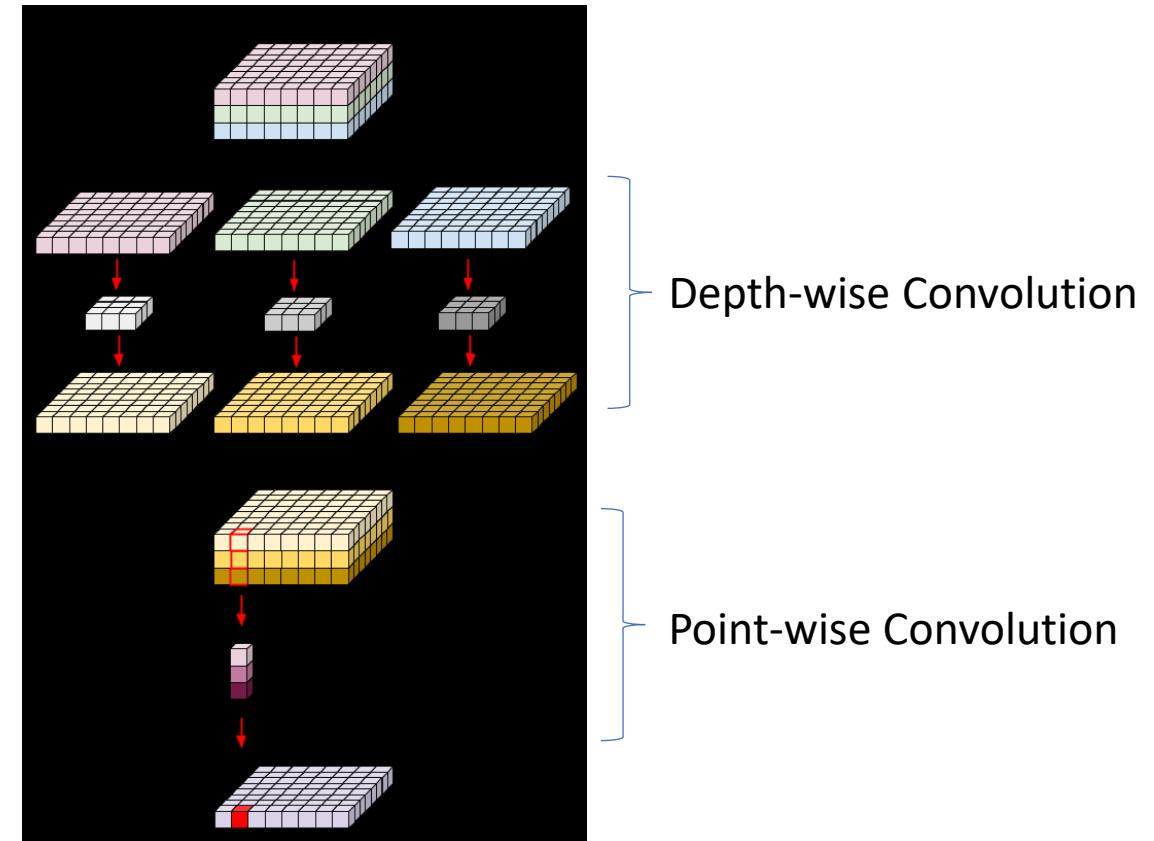
Global Average Pooling as replacement to FC layers:

- An alternative is to use spatial average of feature maps.
- Huge reduction the number of parameters as compared to the Fully Connected layer.
- Stronger local modelling using the micro network.
- It is itself a structural regularizer and hence doesn't need dropout.

Modified Convolution Operation



Standard Convolution

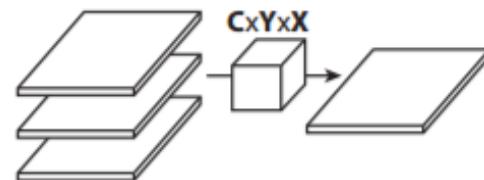


Depth-wise Convolution

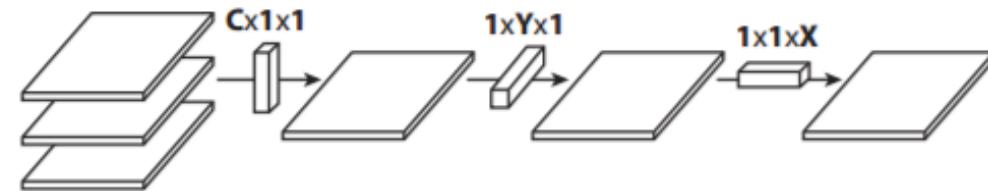
Point-wise Convolution

Flattened Convolutions: Fire Layer

- Replace $c \times y \times x$ convolutions with $c \times 1 \times 1$, $1 \times y \times 1$, and $1 \times 1 \times x$ convolutions



3D convolution

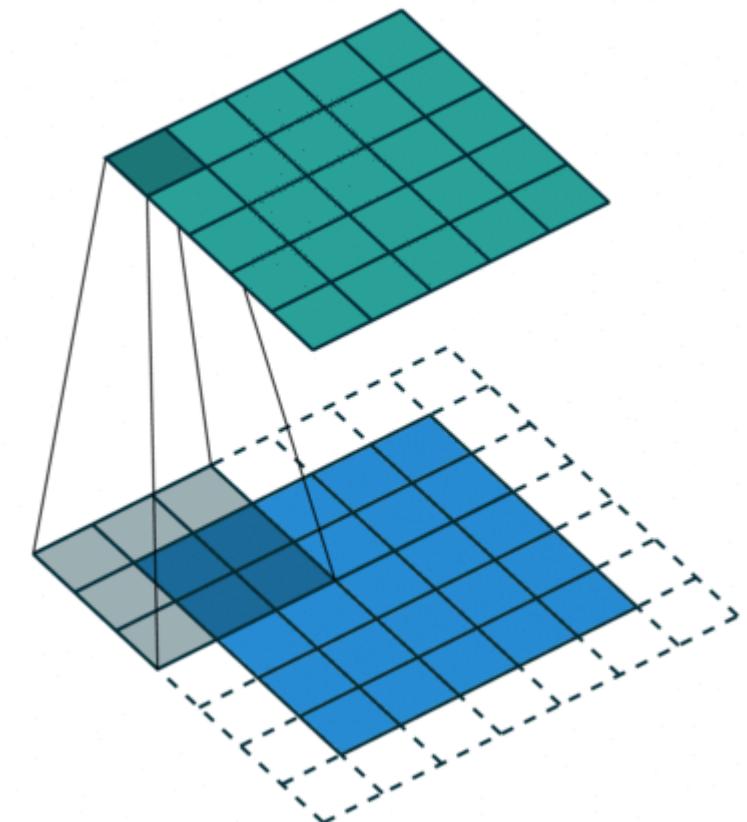


1D convolution over different direction

Iandola, Forrest N., et al. ["SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and< 0.5 MB model size."](#)

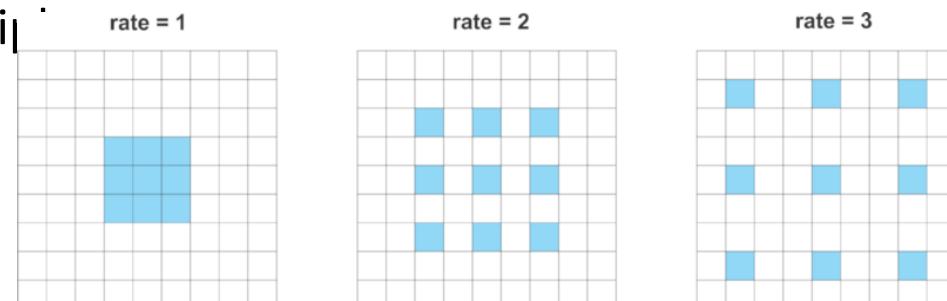
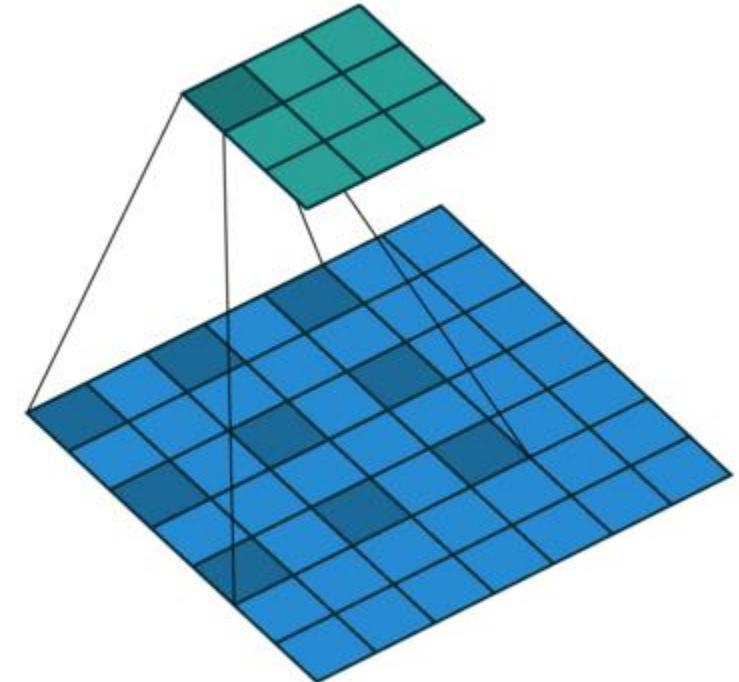
2D Convolution

- **Kernel Size:** The kernel size defines the field of view of the convolution. A common choice for 2D is 3 i.e 3x3 pixels.
- **Stride:** The stride defines the step size of the kernel when traversing the image. While its default is usually 1, we can use a stride of 2 for downsampling an image similar to MaxPooling.
- **Padding:** The padding defines how the border of a sample is handled. A (half) padded convolution will keep the spatial output dimensions equal to the input, whereas unpadded convolutions will crop away some of the borders if the kernel is larger than 1.



Dilated Convolution

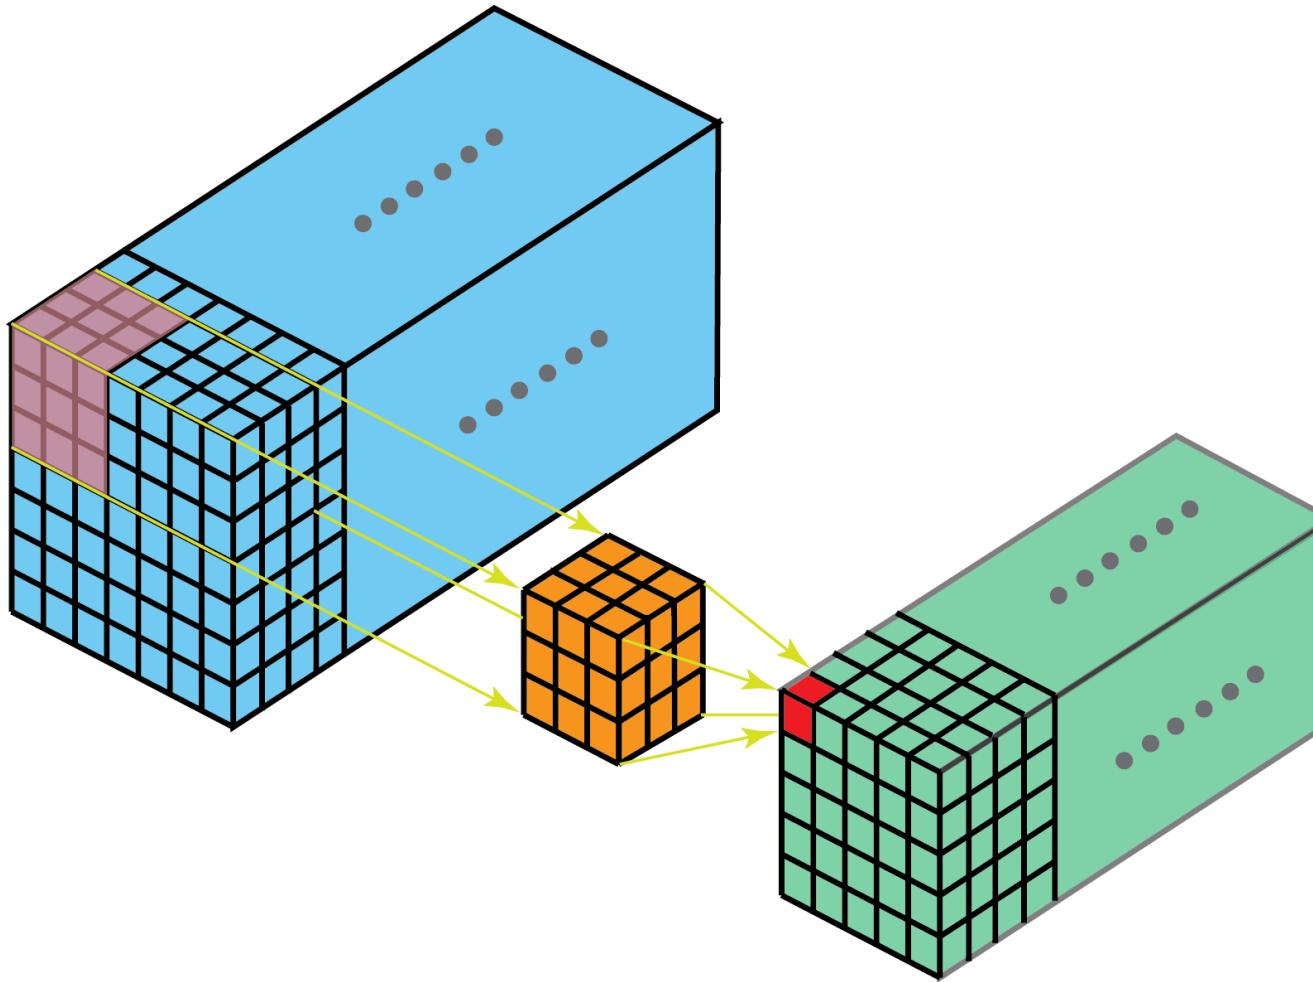
- Dilated convolutions are particularly **popular in the field of real-time segmentation**.
- Dilated convolutions introduce another parameter to convolutional layers called the **dilation rate**, that defines a spacing between the values in a kernel.
- A 3x3 kernel with a dilation rate of 2 will have the same field of view as a 5x5 kernel, while only using 9 parameters. Imagine taking a 5x5 kernel and deleting every second column and row.
- Use them if you need a wide field of view and cannot afford multiple convolutions or larger kernels.



Dilated Convolution

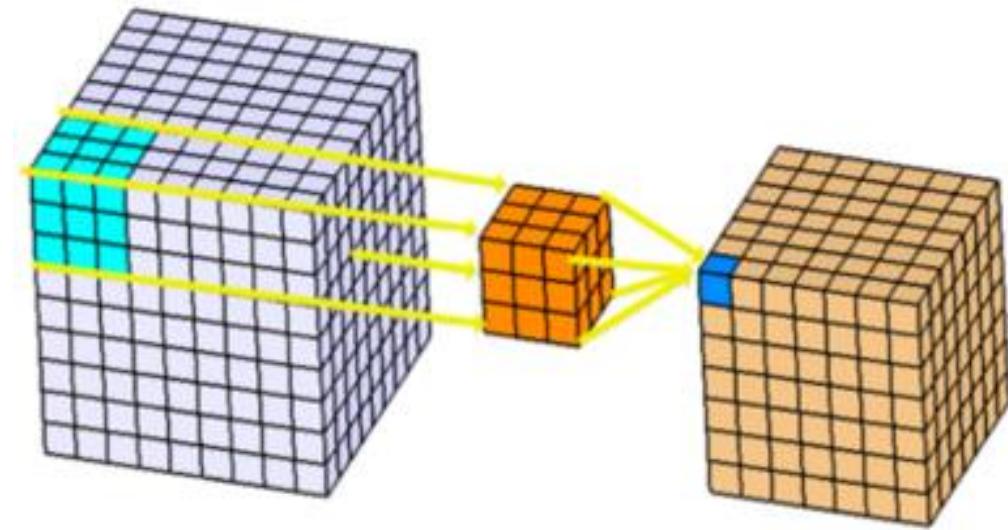
- Dilated convolutions have generally improved performance in semantic segmentation results
- The architecture is based on the fact that dilated convolutions support exponential expansion of the receptive field without loss of resolution or coverage.
- Allows one to have larger receptive field with same computation and memory costs while also preserving resolution.
- Pooling and Strided Convolutions are similar concepts but both reduce the resolution.

3D convolution



0 ₂	0 ₀	0 ₁	0	0	0	0	0
0 ₁	2 ₀	2 ₀	3	3	3	3	0
0 ₀	0 ₁	1 ₁	3	0	3	0	0
0	2	3	0	1	3	0	0
0	3	3	2	1	2	0	0
0	3	3	0	2	3	0	0
0	0	0	0	0	0	0	0

1	6	5
7	10	9
7	10	8

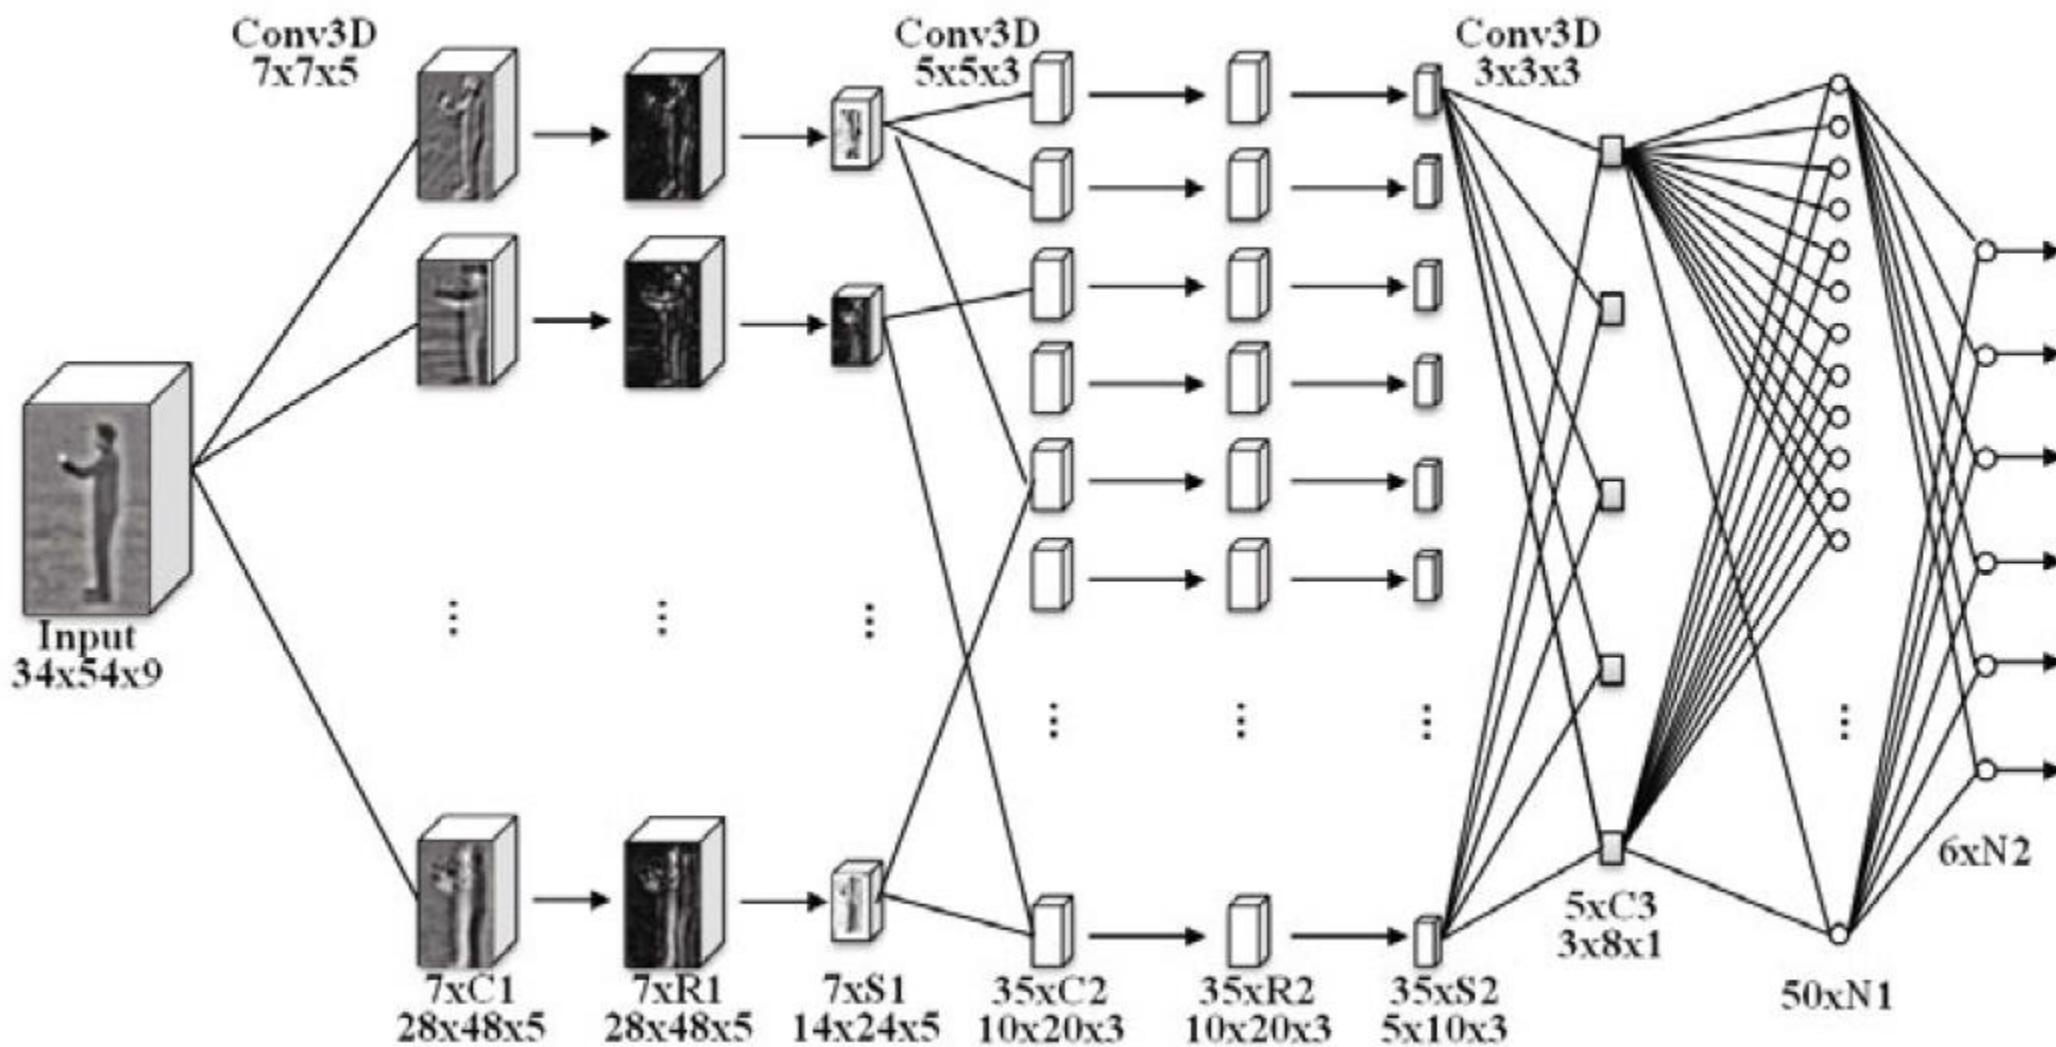


2D vs 3D Convolution

3D CNN

- 3D imageries such as MRI Scan, Human action in video sequence should be processed frame-by-frame
- Temporal property within the 3D imagery is important in extracting the features
- 2D convolution extracts features from only spatial dimensions
- 3D convolution operates on the input volume not only in the X and Y dimensions but also in Z dimension
- Building a 3D CNN requires, 3D Convolution as well as 3D pooling

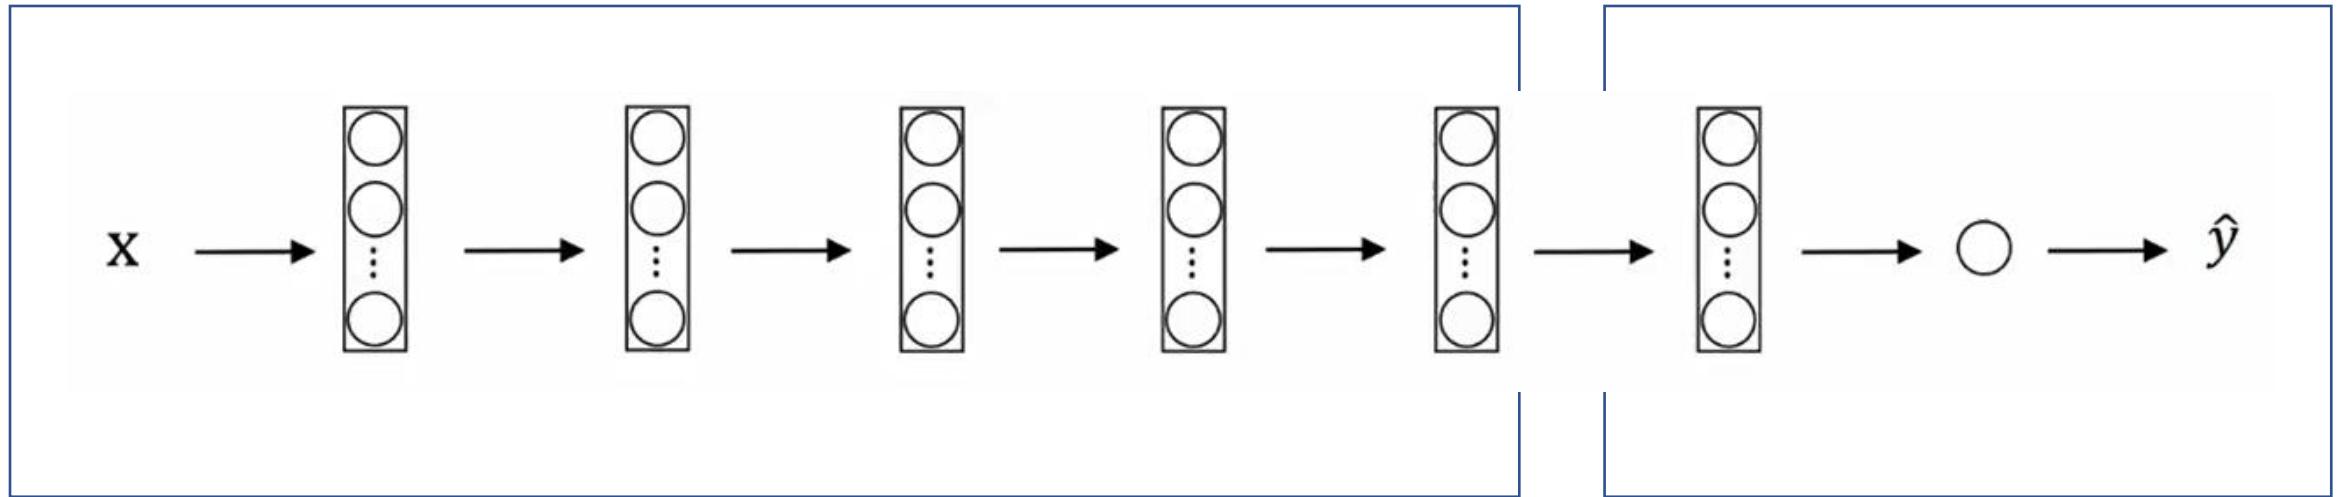
3D CNN for Action Recognition



Transfer Learning

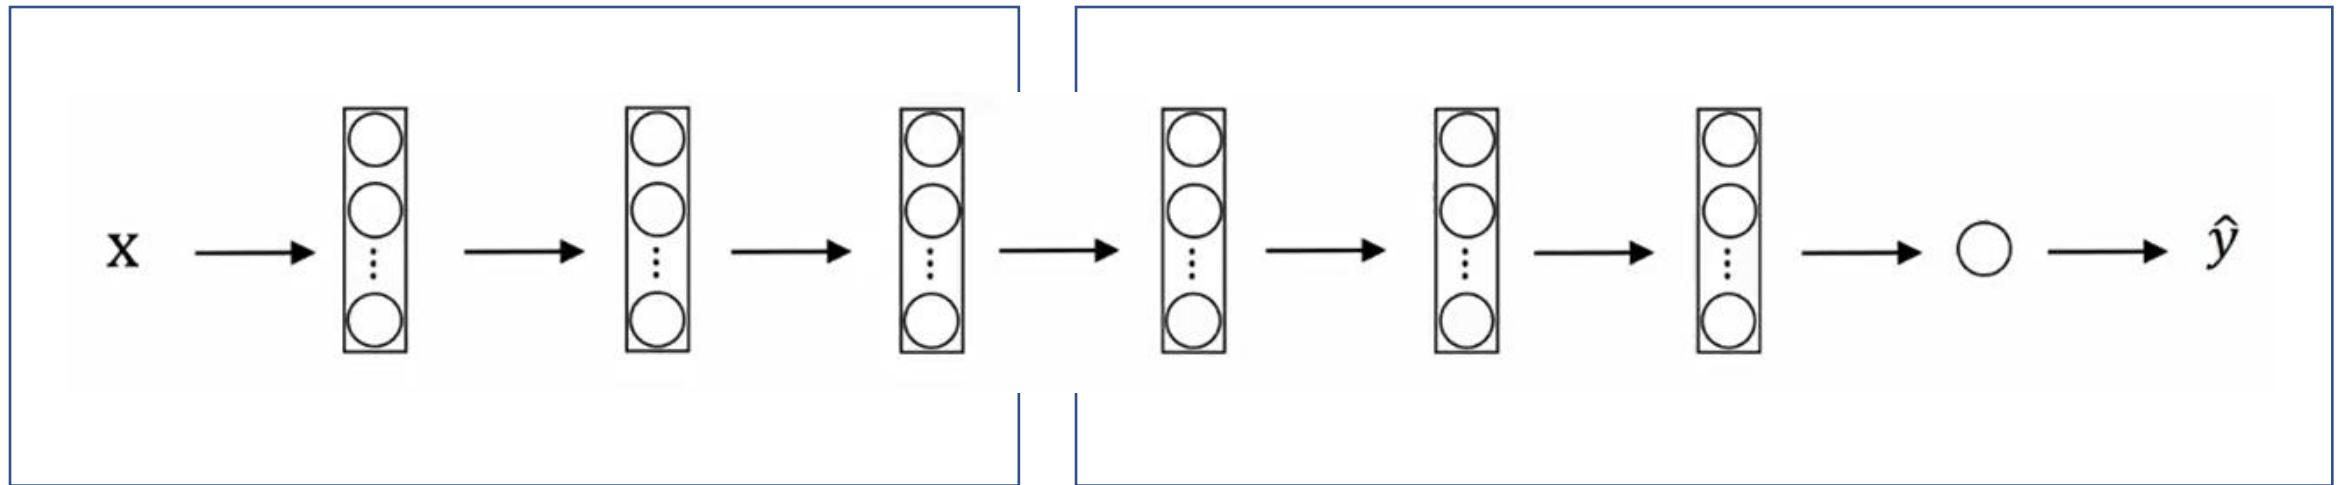
- Sharing the Knowledge gained solving one problem and applying to a different but related problem
- Transfer Learning is next popular driver of deep learning after supervised learning
- The feature spaces of the source and target domain are different, e.g. the documents are written in two different languages.
- The marginal probability distributions of source and target domain are different, e.g. the documents discuss different topics.
- Simulation training is becoming a hot area within the sphere of Deep Learning. Few labs have also started using AR/VR Technologies to be integrated for making advance learning models for some of the critical problems area

Transfer Learning – Small Data



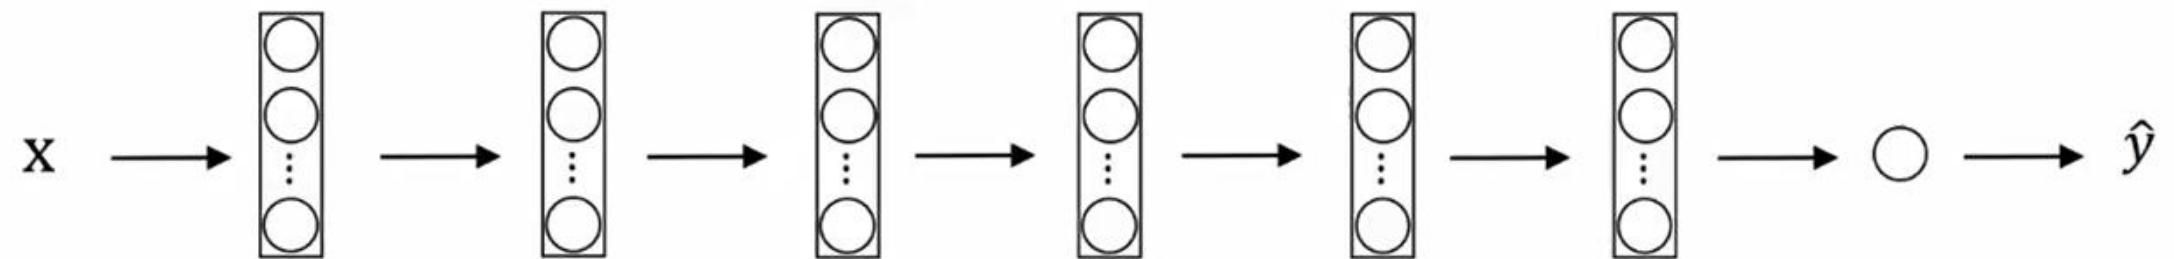
- Finding a Bengal Cat or British Cat where you have small dataset of these categories of cats
- You can download a cat classifier and train last few layers on your data specifically
- In Popular Deep Learning platforms, now you have good support for transfer learning
- Functions like Trainable_Parameters and Freezing specific layers are available

Transfer Learning-Mid Size Data



- In this category we use initial set of layers from the open source model and use the trained weight values.
- Remaining layers there are two ways to handle:
 - we can train the layers from start
 - we can start the weight optimization from the existing set of weights that we have received from the open source model

Transfer Learning- Enough Data



- Where we have enough data to train our new model, open source model may still be useful.
- We can use the pre-trained weights of the model from the same domain and consider that as a starting point for our model.
- It may be much easier and faster to adapt these model for new set of data that we want to train upon.

Face Identification Vs. Verification

- Verification
 - Input an image and also the Name/ID of the person
 - Output whether the input image is of the claimed person
- Identification
 - Input an Image
 - Output whether the image is any of the K persons in the database

Dimension and Parameters

Output dimension and number of parameters

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)			96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)			3x3 pool, stride 2
Conv2	Convolution (5x5)			256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)			3x3 pool, stride 2
Conv3	Convolution (3x3)			384 filters, stride 1, padding 1
Conv4	Convolution (3x3)			384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

Output dimension and number of parameters (Conv1)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54 , 54 , 96)	$11 * 11 * 3 * 96 + 96$	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26 , 26 , 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(25 , 25 , 256)	$5 * 5 * 96 * 256 + 256$	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12 , 12 , 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12 , 12 , 384)	$3 * 3 * 256 * 384 + 384$	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)			384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters = (Filter_Height × Filter_Width × No. of Input Channels) × No. of Filters + Number of Filters

$$\text{No. of Parameters} = (11 * 11 * 3 * 96) + 96 = 34,944$$

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(224 - 11 + 2 * 0)}{4} + 1 = \frac{213}{4} + 1 = 54$$

Output dimension and number of parameters (**Pool1**)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)			3x3 pool, stride 2
Conv2	Convolution (5x5)			256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)			3x3 pool, stride 2
Conv3	Convolution (3x3)			384 filters, stride 1, padding 1
Conv4	Convolution (3x3)			384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(54 - 3 + 2 * 0)}{2} + 1 = 26$$

Output dimension and number of parameters (**Conv2**)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)			256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)			3x3 pool, stride 2
Conv3	Convolution (3x3)			384 filters, stride 1, padding 1
Conv4	Convolution (3x3)			384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters = (Filter_Height × Filter_Width × No. of Input Channels) × No. of Filters + Number of Filters

$$\text{No. of Parameters} = (5 * 5 * 96 * 256) + 256 = 6,14,656$$

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(26 - 5 + 2 * 2)}{1} + 1 = 26$$

Output dimension and number of parameters (**Pool2**)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)			3x3 pool, stride 2
Conv3	Convolution (3x3)			384 filters, stride 1, padding 1
Conv4	Convolution (3x3)			384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(26 - 3 + 2 * 0)}{2} + 1 = 12$$

Output dimension and number of parameters (Conv3)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)			384 filters, stride 1, padding 1
Conv4	Convolution (3x3)			384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters = (Filter_Height × Filter_Width × No. of Input Channels) × No. of Filters + Number of Filters

$$\text{No. of Parameters} = (3 * 3 * 256 * 384) + 384 = 8,85,120$$

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(12 - 3 + 2 * 1)}{1} + 1 = 12$$

Output dimension and number of parameters (Conv4)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	8,85,120	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)			384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters = (Filter_Height × Filter_Width × No. of Input Channels) × No. of Filters + Number of Filters

$$\text{No. of Parameters} = (3 * 3 * 384 * 384) + 384 = 13,27,488$$

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(12 - 3 + 2 * 1)}{1} + 1 = 12$$

Output dimension and number of parameters (Conv5)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	8,85,120	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)	(12, 12, 384)	13,27,488	384 filters, stride 1, padding 1
Conv5	Convolution (3x3)			256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters = (Filter_Height × Filter_Width × No. of Input Channels) × No. of Filters + Number of Filters

$$\text{No. of Parameters} = (3 * 3 * 384 * 256) + 256 = 8,84,992$$

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(12 - 3 + 2 * 1)}{1} + 1 = 12$$

Output dimension and number of parameters (**Pool3**)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	8,85,120	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)	(12, 12, 384)	13,27,488	384 filters, stride 1, padding 1
Conv5	Convolution (3x3)	(12, 12, 256)	8,84,992	256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)			3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

$$\text{Output Size} = \frac{(\text{Input Size} - \text{Filter Size} + 2 * \text{Padding})}{\text{Stride}} + 1$$

$$\text{Output Size} = \frac{(12 - 3 + 2 * 0)}{2} + 1 = 5$$

Output dimension and number of parameters (**Flatten**)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	8,85,120	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)	(12, 12, 384)	13,27,488	384 filters, stride 1, padding 1
Conv5	Convolution (3x3)	(12, 12, 256)	8,84,992	256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)	(5, 5, 256)	0	3x3 pool, stride 2
Flatten	-			Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

$$\text{Flatten} = \text{Height} * \text{Width} * \text{Channel} = 5 * 5 * 256 = 6400$$

Output dimension and number of parameters (FC1)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	8,85,120	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)	(12, 12, 384)	13,27,488	384 filters, stride 1, padding 1
Conv5	Convolution (3x3)	(12, 12, 256)	8,84,992	256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)	(5, 5, 256)	0	3x3 pool, stride 2
Flatten	-	(6400 X 1)	0	Flatten to 1D
FC1	Fully Connected			4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters = $(6400 * 4096) + 4096 = 2,62,14,656$

Output dimension and number of parameters (FC2)

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	8,85,120	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)	(12, 12, 384)	13,27,488	384 filters, stride 1, padding 1
Conv5	Convolution (3x3)	(12, 12, 256)	8,84,992	256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)	(5, 5, 256)	0	3x3 pool, stride 2
Flatten	-	(6400 X 1)	0	Flatten to 1D
FC1	Fully Connected	4096	2,62,14,656	4096 units
FC2	Fully Connected			4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters=(4096 * 4096) + 4096 = 1,67,81,312

Output dimension and number of parameters (FC3)

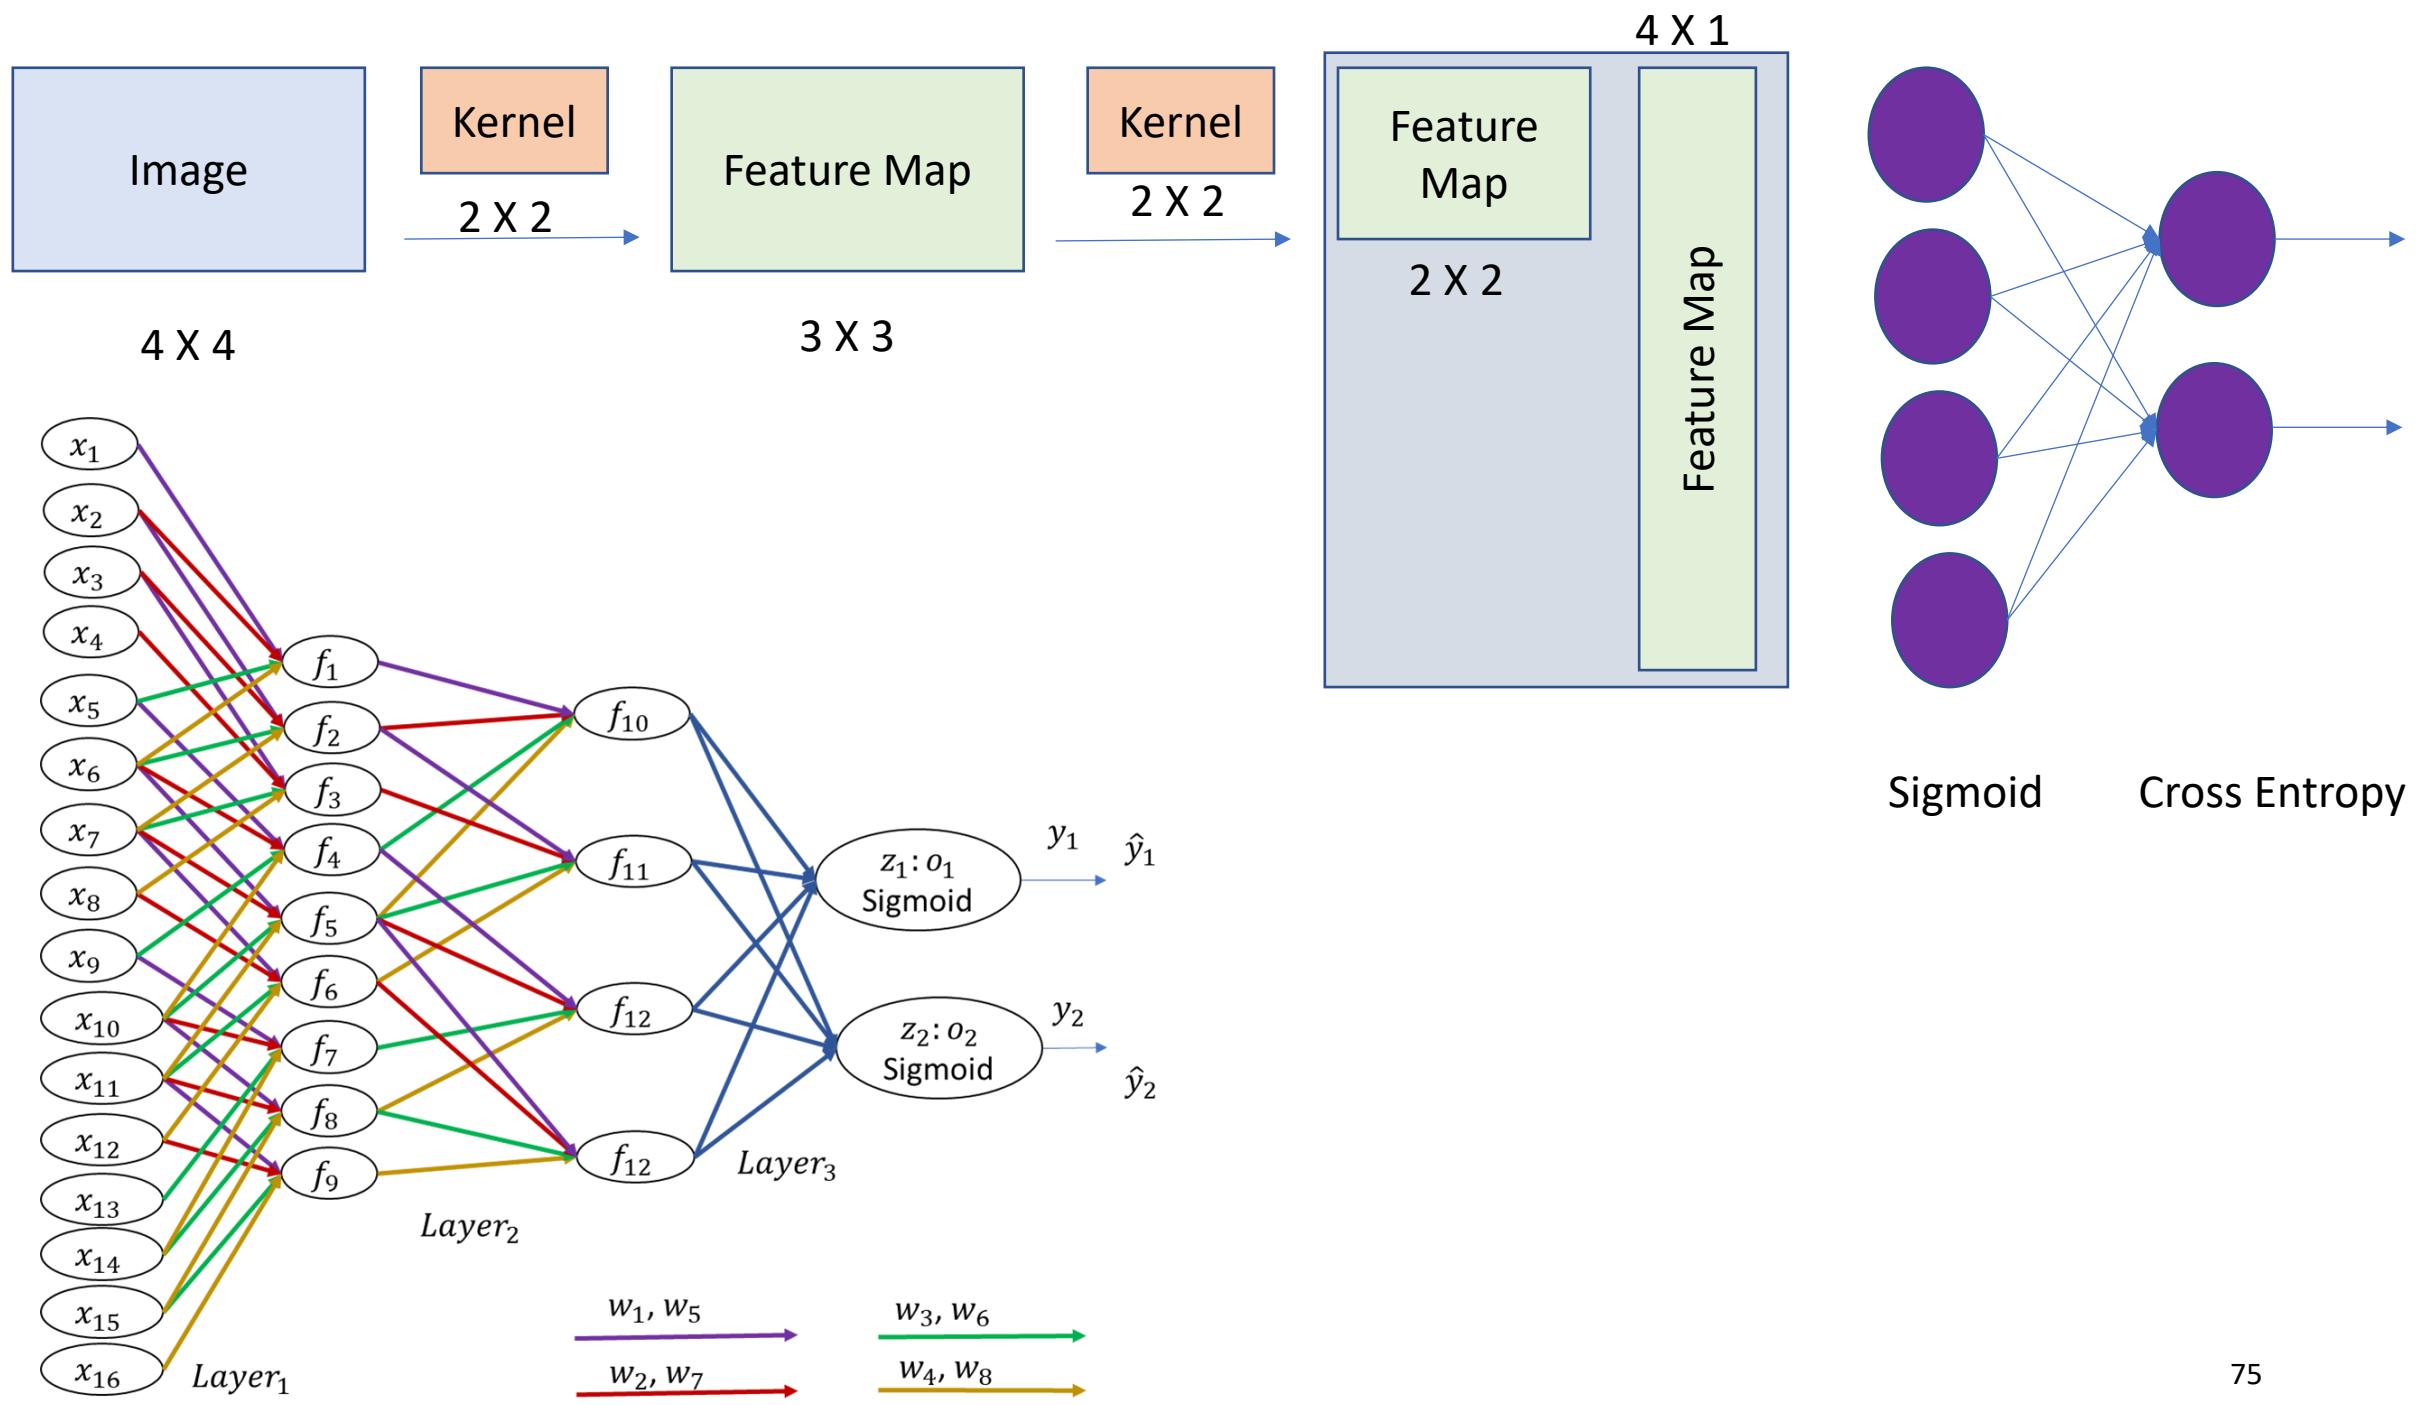
Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	34,944	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	6,14,656	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	8,85,120	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)	(12, 12, 384)	13,27,488	384 filters, stride 1, padding 1
Conv5	Convolution (3x3)	(12, 12, 256)	8,84,992	256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)	(5, 5, 256)	0	3x3 pool, stride 2
Flatten	-	(6400 X 1)	0	Flatten to 1D
FC1	Fully Connected	4096	2,62,14,656	4096 units
FC2	Fully Connected	4096	1,67,81,312	4096 units
FC3 (Output)	Fully Connected			1000 units for classification

No. of Parameters = $(4096 * 1000) + 1000 = 40,97,000$

Output dimension and number of parameters

Layer	Type	Output Shape	Number of Parameters	Explanation
Input	-	(224, 224, 3)	0	Input RGB image
Conv1	Convolution (11x11)	(54, 54, 96)	$(11*11*3*96) + 96 = 34,944$	96 filters, stride 4, padding 0
Pool1	Max Pooling (3x3)	(26, 26, 96)	0	3x3 pool, stride 2
Conv2	Convolution (5x5)	(26, 26, 256)	$(5*5*96*256) + 256 = 614,656$	256 filters, stride 1, padding 2
Pool2	Max Pooling (3x3)	(12, 12, 256)	0	3x3 pool, stride 2
Conv3	Convolution (3x3)	(12, 12, 384)	$(3*3*256*384) + 384 = 885,120$	384 filters, stride 1, padding 1
Conv4	Convolution (3x3)	(12, 12, 384)	$(3*3*384*384) + 384 = 1,327,488$	384 filters, stride 1, padding 1
Conv5	Convolution (3x3)	(12, 12, 256)	$(3*3*384*256) + 256 = 884,992$	256 filters, stride 1, padding 1
Pool3	Max Pooling (3x3)	(5, 5, 256)	0	3x3 pool, stride 2
Flatten	-	6400	0	Flatten to 1D
FC1	Fully Connected	4096	$(6400*4096) + 4096 = 26,214,656$	4096 units
FC2	Fully Connected	4096	$(4096*4096) + 4096 = 16,781,312$	4096 units
FC3 (Output)	Fully Connected	1000	$(4096*1000) + 1000 = 4,097,000$	1000 units for classification

Backpropogation CNN



Thank You